AZURE DATA FACTORY – DATA FLOWS

BRINGING THE "T" IN ELT

WHO AM I?

t: @bizdataviz
In/cseferlis

- Practice Manager – Pragmatic Works – www.pragmaticworks.com
- Former CIO
- Blog, Speak, Record, Tweet
  - Twitter: @bizdataviz
  - LinkedIn: in/cseferlis
  - http://blog.bizdataviz.com
- MBA from UMass
- Ski, Bike, Hike, Run and drag the family

along…

WHY ARE WE HERE?

t: @bizdataviz
ln/cseferlis

- Modern Data Warehouse (Platform)
- Azure Data Orchestration
- Preview New Azure Data Factory Features
- Full Circle ELT with Azure Data Factory

AZURE DATA FACTORY… QUICKLY

t: @bizdataviz
In/cseferlis

- Native Cloud Data Orchestration Tool
- Started with v1, but lacked many features
- v2 goes GA 6/2018 – work has begun on Data Flows
- Data Flows and other new features currently in Limited Preview

LIMITED PREVIEW FEATURES

t: @bizdataviz
In/cseferlis

- Debug Mode
- Data Inspection
- Transformations
- GitHub Support
- Detailed Flow Inspection
- All-New Expression Builder

VISUAL DATA FLOW AUTHORING

t: @bizdataviz
In/cseferlis

- Transform Data, At Scale, in the Cloud, Zero-Code
  - Cloud-first, scale-out ELT
  - Code-free dataflow pipelines
- Serverless scale-out transformation execution engine
- Maximum Productivity for Data Engineers
  - Does NOT require understanding of Spark / Scala / Python / Java
- Resilient Data Transformation Flows
  - Built for big data scenarios with unstructured data requirements
  - Operationalize with Data Factory scheduling, control flow and monitoring

CODE-FREE DATA TRANSFORMATION AT SCALE

t: @bizdataviz
In/cseferlis

- Does not require understanding of Spark, Big Data Execution Engines, Clusters, Scala, Python …
- Focus on building business logic and data transformation
  - Data cleansing
  - Aggregation
  - Data conversions
  - Data prep
  - Data exploration

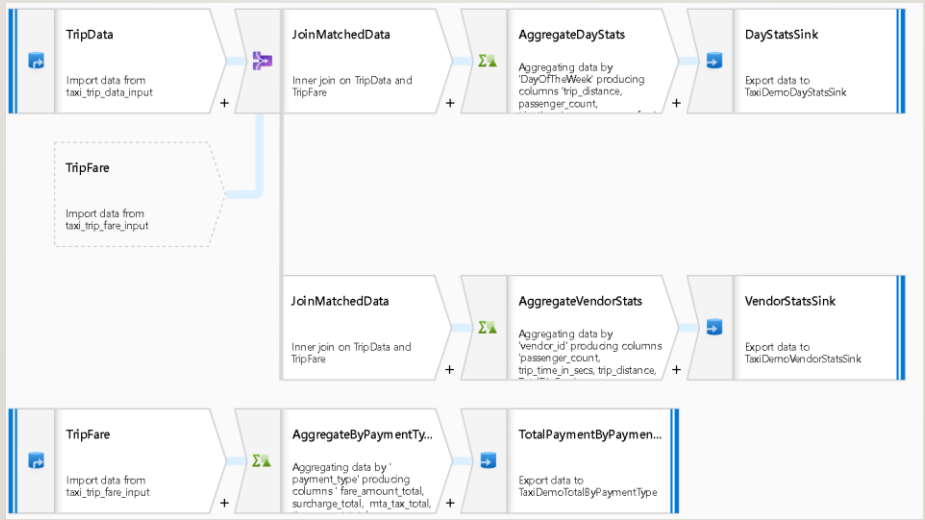

… not …

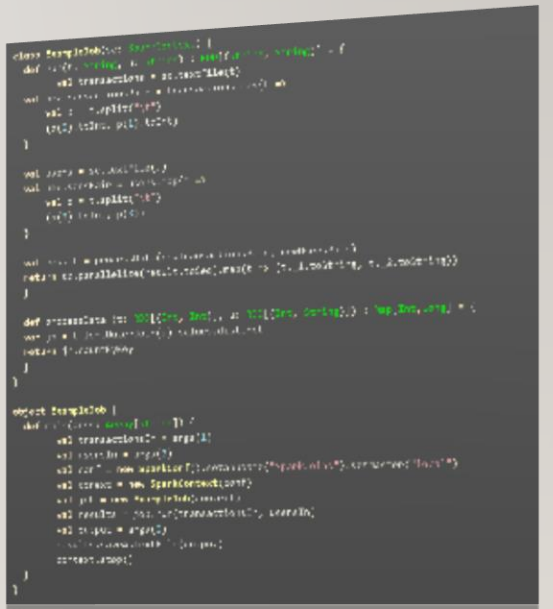

ADF DATA FLOW WORKSTREAM

t: @bizdataviz
In/cseferlis

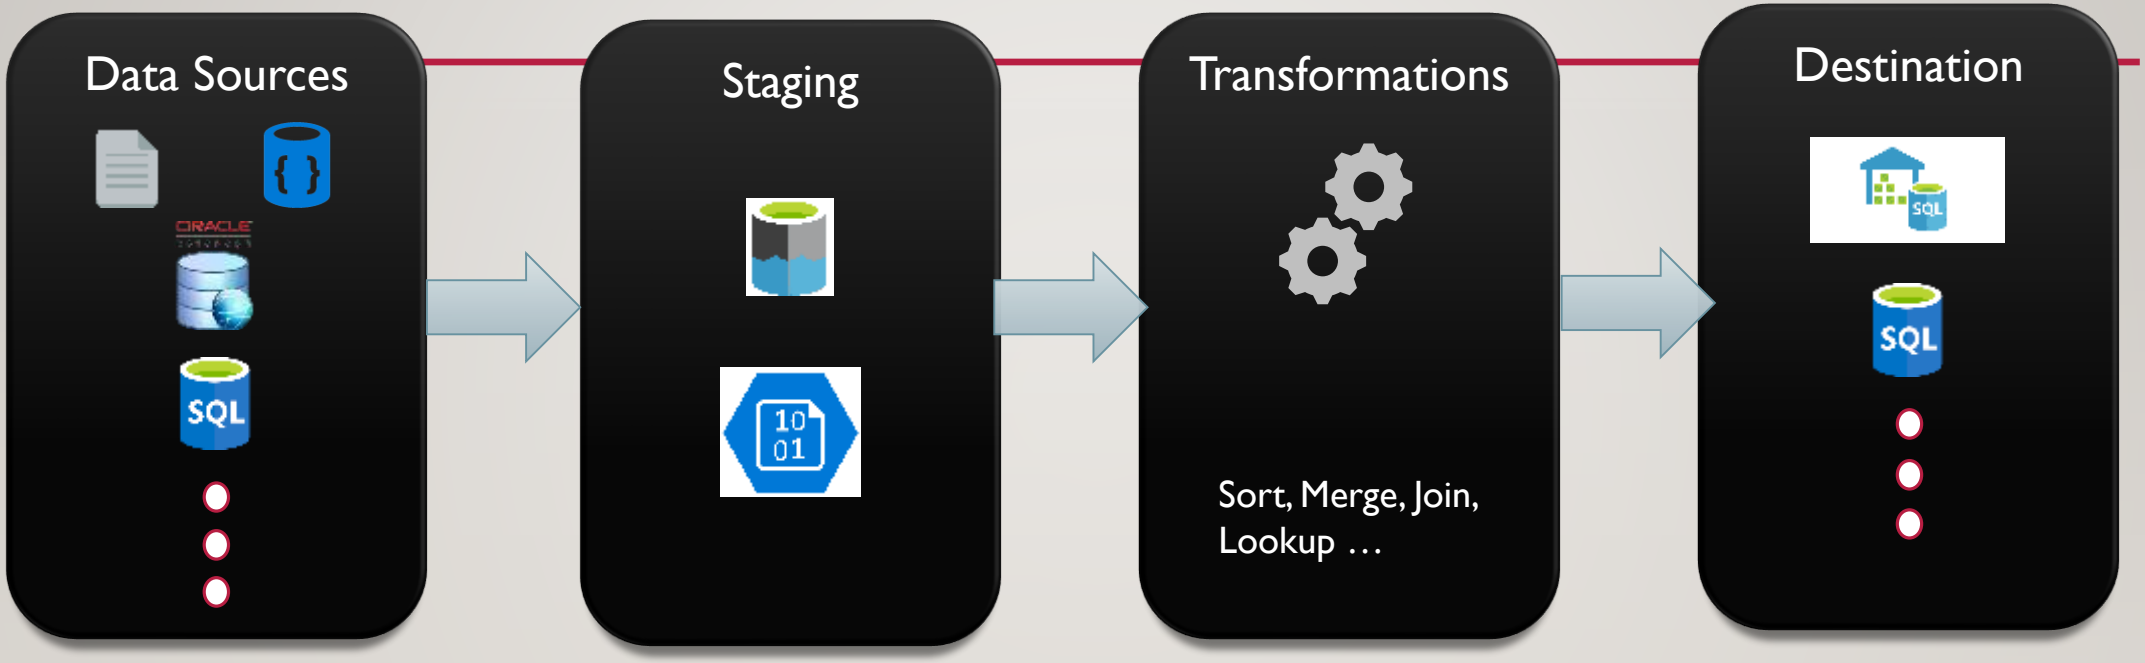

- Explicit user action
- User places data source(s) on design surface, from toolbox
- Select explicit sources

- Implicit/Explicit
- Data Lake staging area as default
- User does not need to configure this manually

- Explicit user action
- User places transformations on design surface, from toolbox
- User must set properties for transformation steps and step connectors

- Explicit user action
- User chooses destination connector(s)
- User sets connector property options

DATA FLOW LIMITED PREVIEW

t: @bizdataviz
In/cseferlis

- Azure SLAs are NA for preview services (private or public preview) until GA of the service.
- Limited Preview Support
  - Handled directly with the Azure Engineering team via adfdataflowext@microsoft.com.
- Sign-up for ADF Data Flow service
  - http://aka.ms/dataflowpreview
  - Microsoft Azure must whitelist your subscription ID to turn on the feature for you
- Public Preview Support
  - Normal Azure customer service channels

t: @bizdataviz
ln/cseferlis

COMMON ELT SCENARIOS HANDLED WITH ADF – DATA FLOWS

SIMPLE COPY FLOW

t: @bizdataviz
In/cseferlis

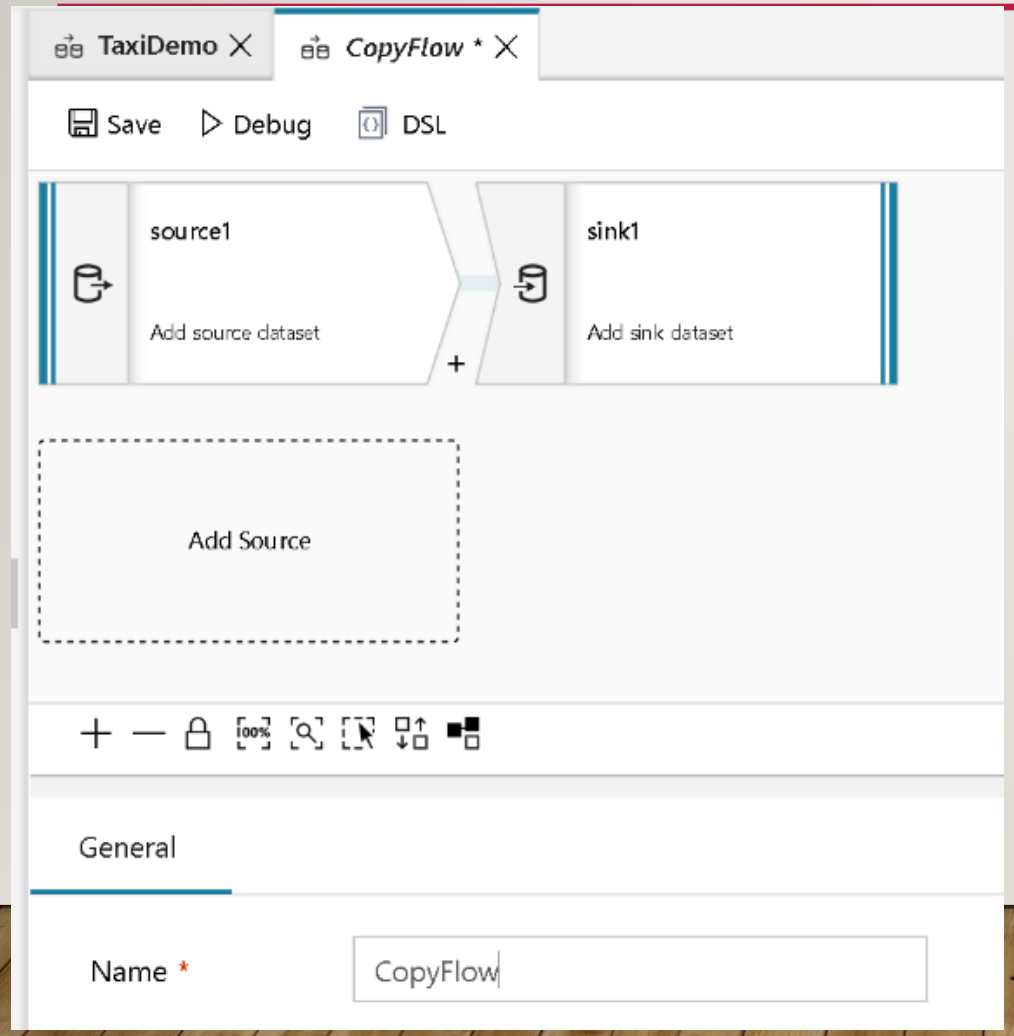

- ADF Data Flow is a guided construction process
- Begin by defining the Datasets for your Source and Sink
- Add Transformations to each node in your data flow
- Or simply copy from source to sink with no transformation
- Map columns and fields along the way

SLOWLY CHANGING DIMENSION

t: @bizdataviz
In/cseferlis

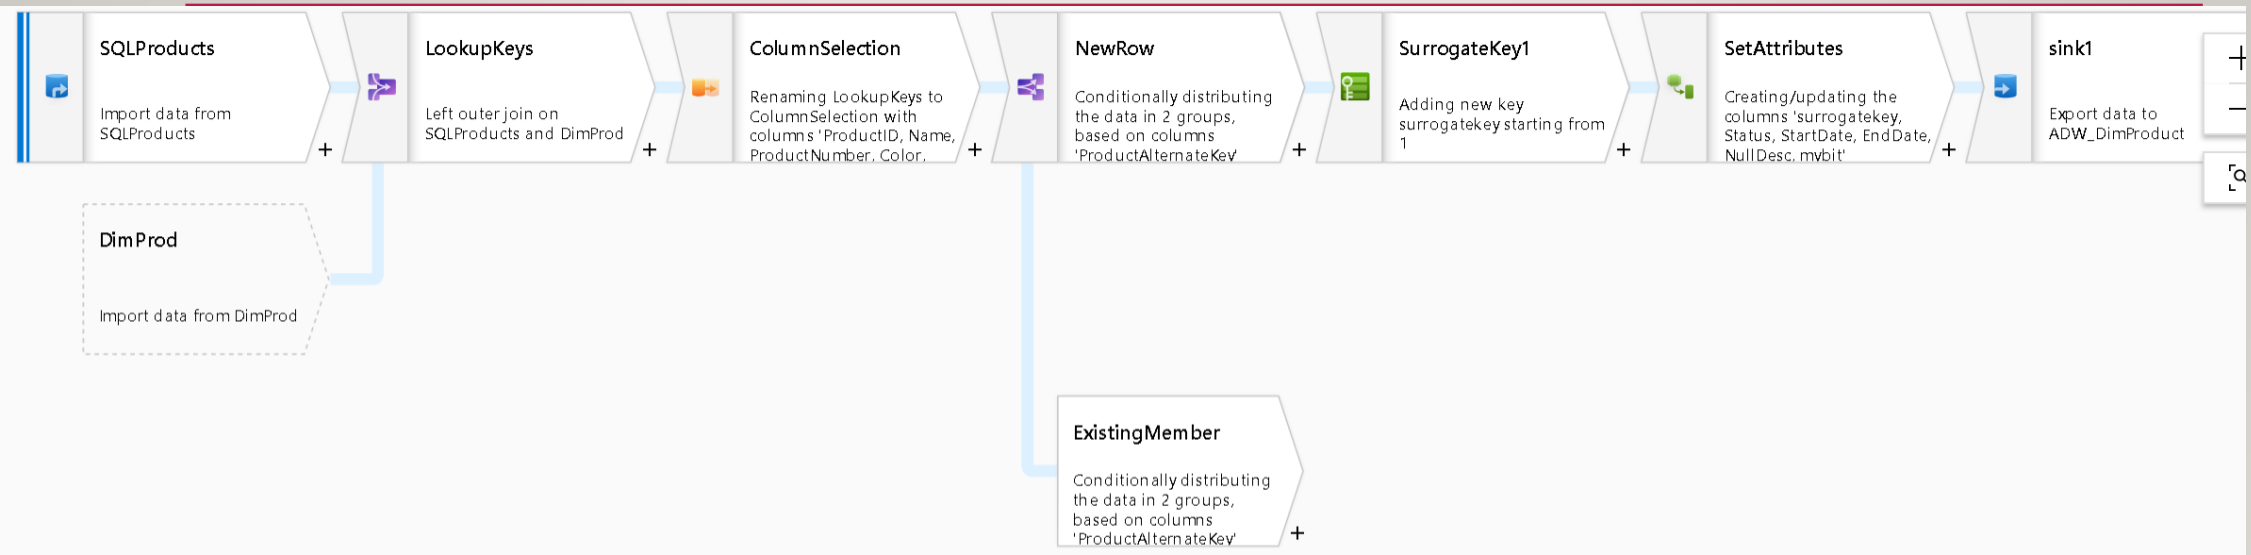

- Common DW pattern to manage changing attributes to dimension members
- Graphically build code-free SCD ETL pattern to load your data warehouse
- Connect directly to Azure SQL DB and Azure SQL DW
- Use Lookup, Surrogate Key, Derived Column and Select transforms

LOAD STAR SCHEMA DW

t: @bizdataviz
In/cseferlis

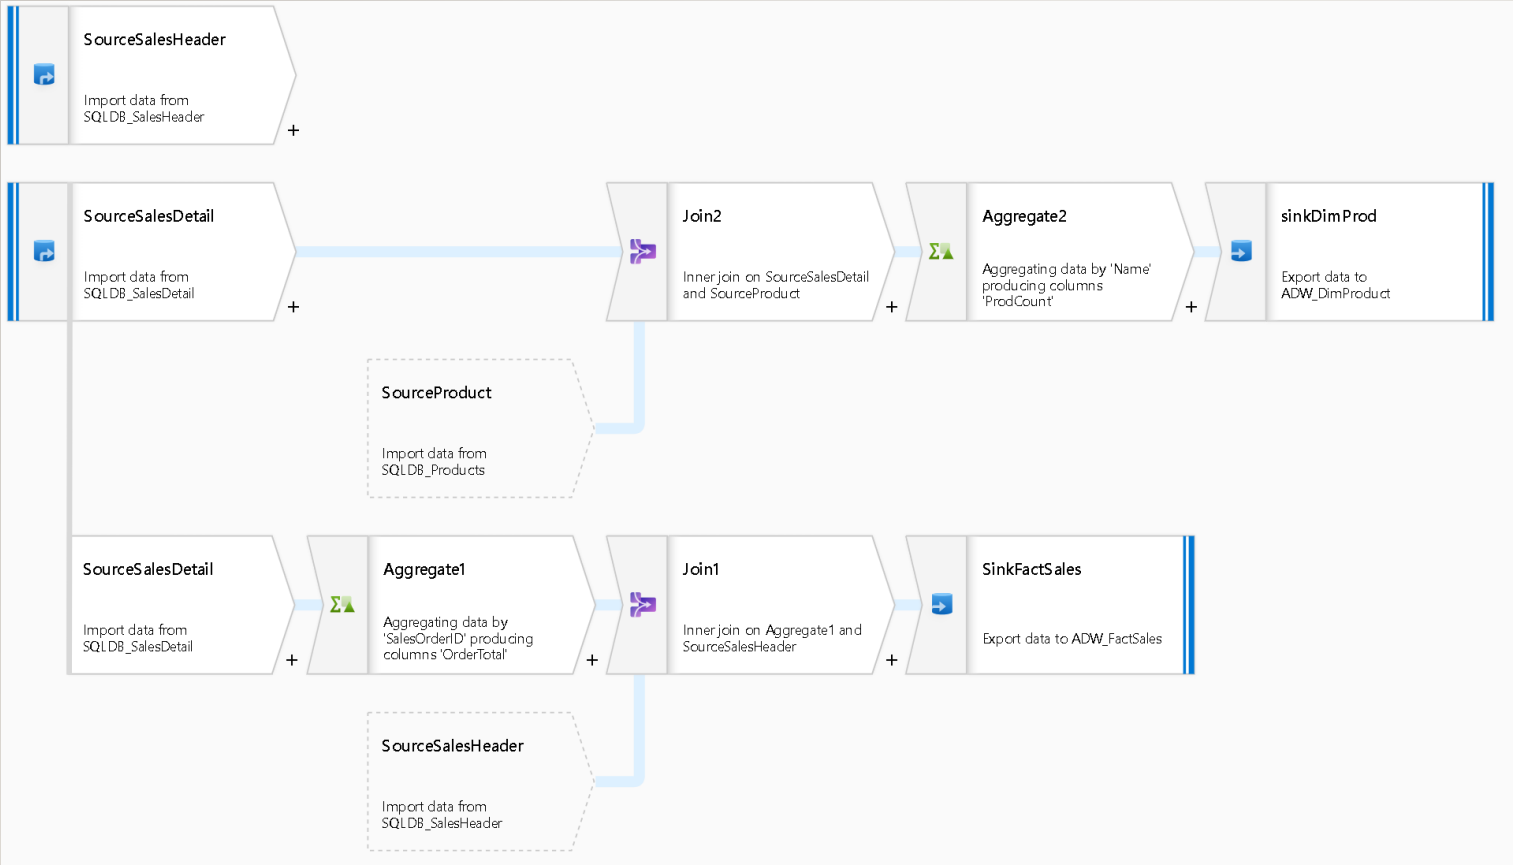

- Classic ETL pattern is easy to build in ADF's code-free Data Flow visual data transformation environment
- Add Aggregate transforms to produce calculations that you store in your analytical database schema
- Use Join transform to combine data from multiple data sources and data streams inside your data flow
- Land your data in your Lake folders or direct to Azure SQL DW

DATA LAKE/DATA SCIENCE

t: @bizdataviz
In/cseferlis

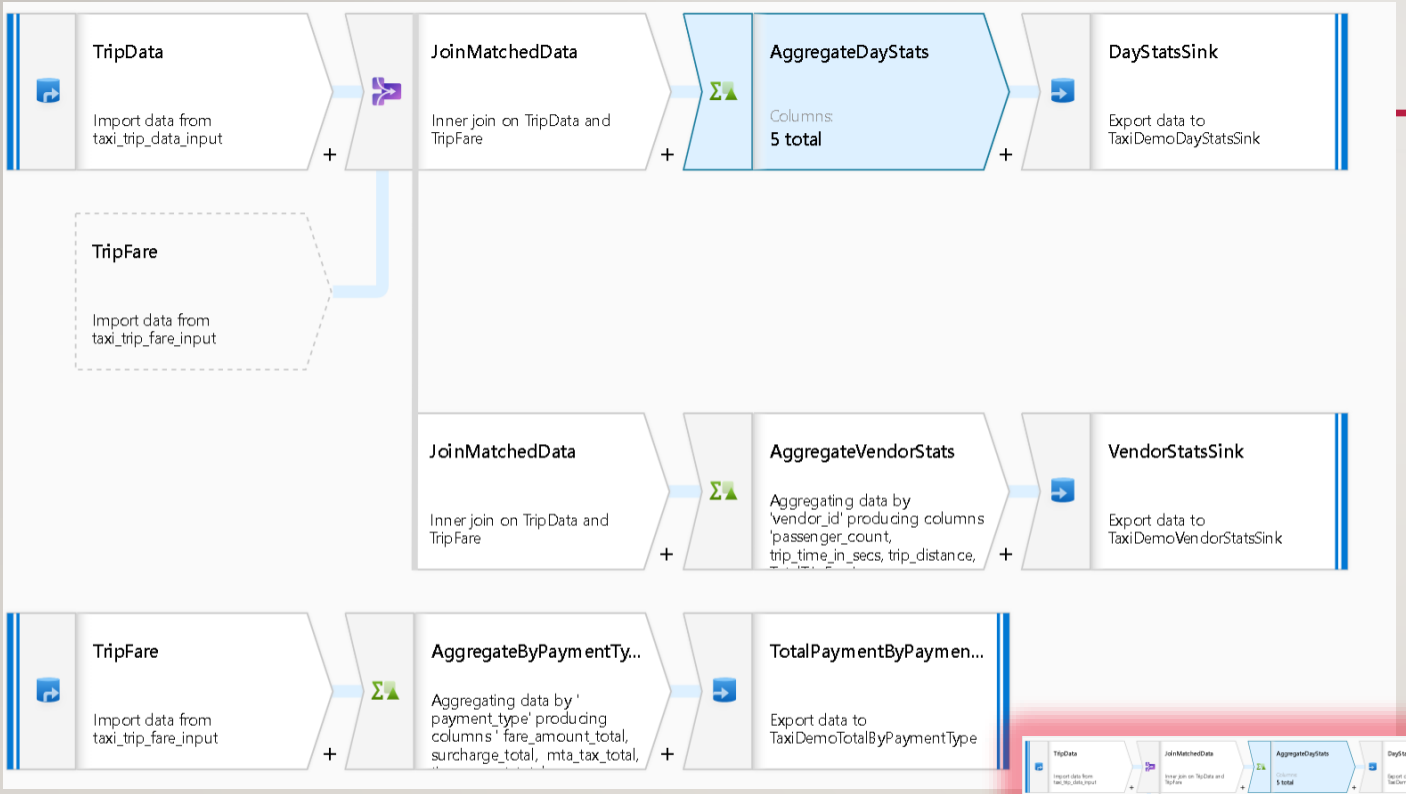

- ADF supports building visual data transformations against your data directly in Data Lake locations (i.e. Azure Blob Store, Azure Data Lake Store)
- Built-in handling of schema drift for frequent changes in data lake file formats, columns, and data types
- Perform data exploration and data profiling across your data lake in ADF Data Flow with interactive debug data preview

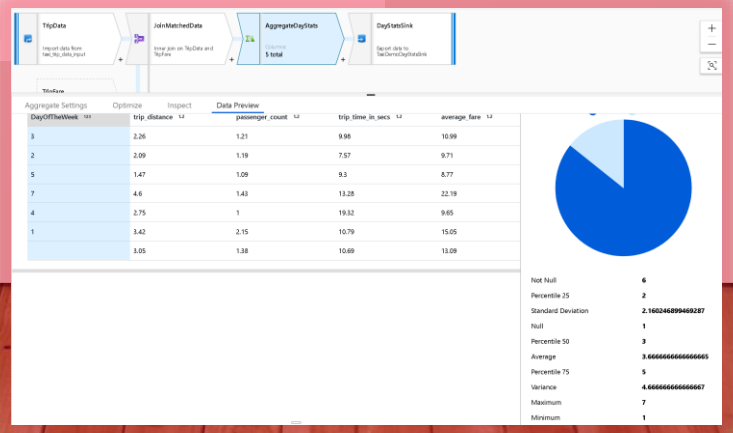

t: @bizdataviz
In/cseferlis

NEW FEATURES IN AZURE DATA FACTORY

INTERACTIVE EXPRESSION BUILDER – BUILD TRANSFORM EXPRESSIONS, NOT CODE

t: @bizdataviz
ln/cseferlis

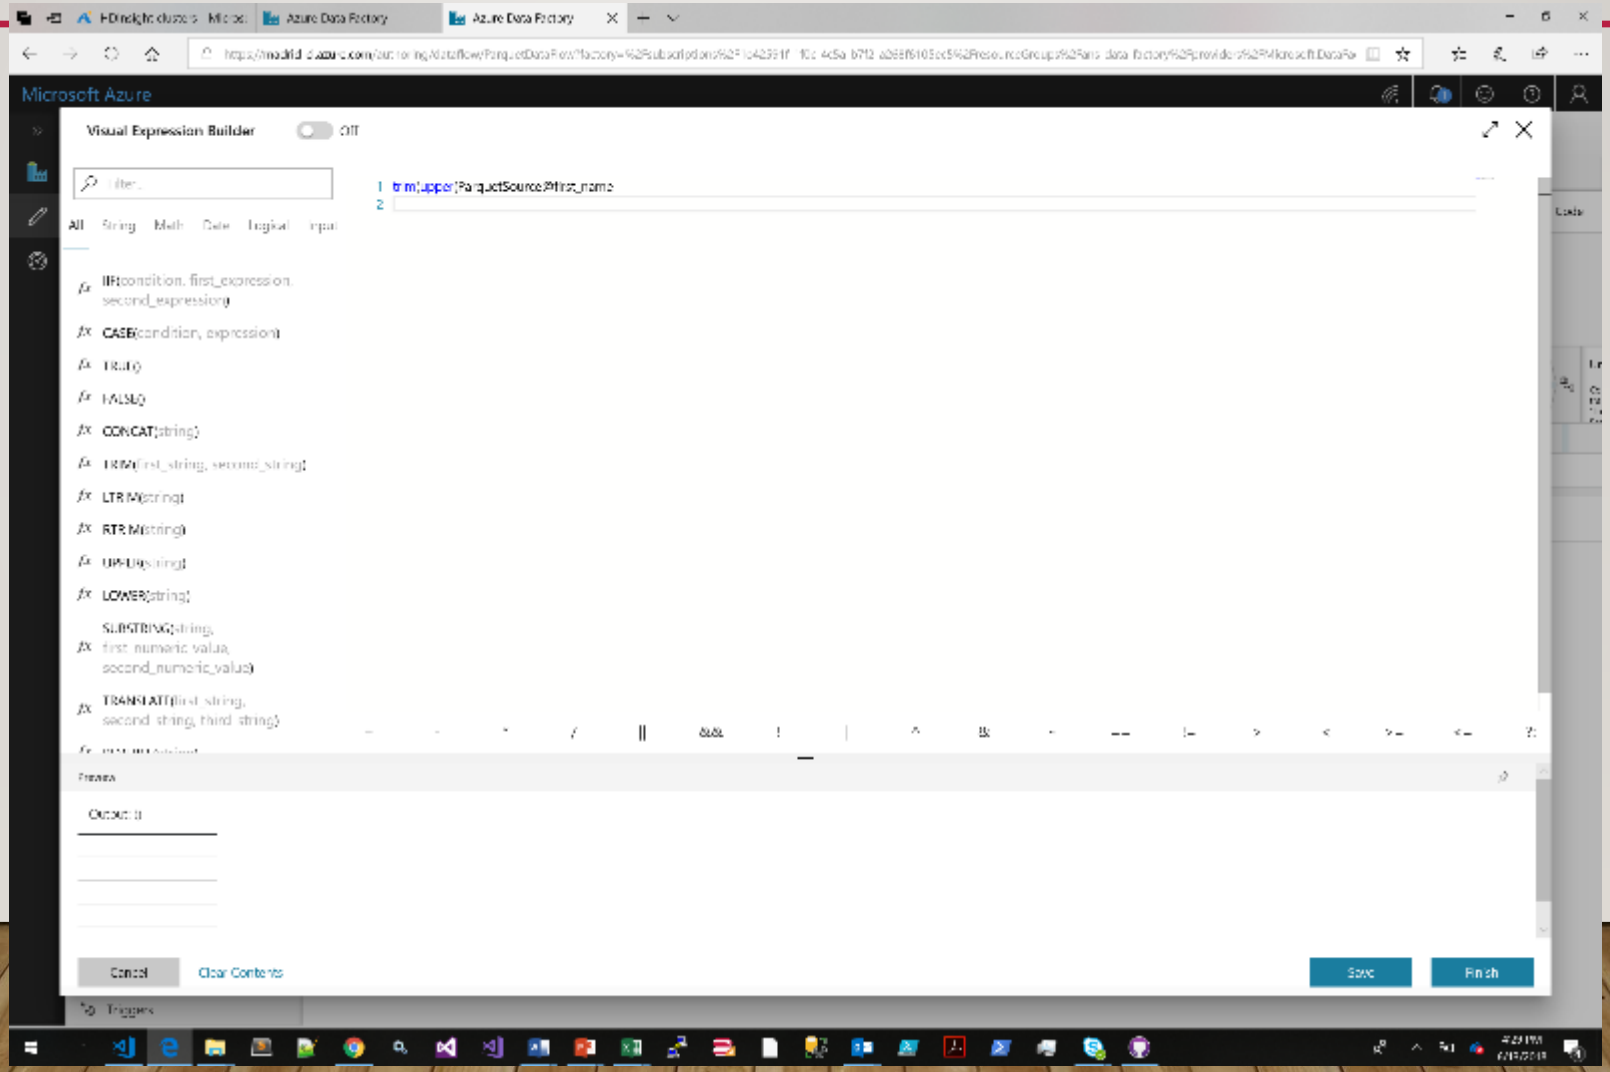

DEBUG DATA FLOWS WITH DATA PREVIEW AND DATA SAMPLING

t: @bizdataviz
In/cseferlis

	Update	New	Unchanged	Total
Nuber of columns	2	1	15	**18**
Nuber of rows	30	0	2,483	**2,234**

Output schema | Data Preview

Order	Date	Column	Profit 1.2	Column 123	Continent_data abc	Column abc
1	12/03/2018	✓	2455.45	Cell Contents	North America	Cell Contents
2	12/03/2018	✓	2455.45	Cell Contents	Europe	Cell Contents
3	12/03/2018		2455.45	Cell Contents	North America	Cell Contents
4	12/03/2018		2455.45	Cell Contents	North America	Cell Contents
5	12/03/2018	✓	2455.45	Cell Contents	Africa	Cell Contents
6	12/03/2018	✓	2455.45	Cell Contents	Asia	Cell Contents
7	12/03/2018	✓	null	Cell Contents	Europe	Cell Contents
8	12/03/2018		null	Cell Contents	North America	Cell Contents
9	12/03/2018		49582.23	Cell Contents	Asia	Cell Contents
10	12/03/2018		49582.23	Cell Contents	Asia	Cell Contents
11	12/03/2018	✓	null	Cell Contents	Europe	Cell Contents
12	12/03/2018		49582.23	Cell Contents	Europe	Cell Contents
13	12/03/2018		49582.23	Cell Contents	North America	Cell Contents

Continent_data abc

528	51.8%	North America
203	20.3	Asia
136	13.6%	Africa
116	11.6%	Europe
17	1.7%	America
10	1.0%	Other

Content Type	String
Missing Value	0 (0%)
Unique Values	5 (5%)
Most Common	North America (52.8%)
Skewness	-0.043210436

t: @bizdataviz
In/cseferlis

BUILD RESILIENT DATA FLOWS WITH SCHEMA DRIFT HANDLING

HANDLING SOURCE CHANGES

t: @bizdataviz
In/cseferlis

- Data Engineer Defines Source and you take ALL fields from source w/flexible schema

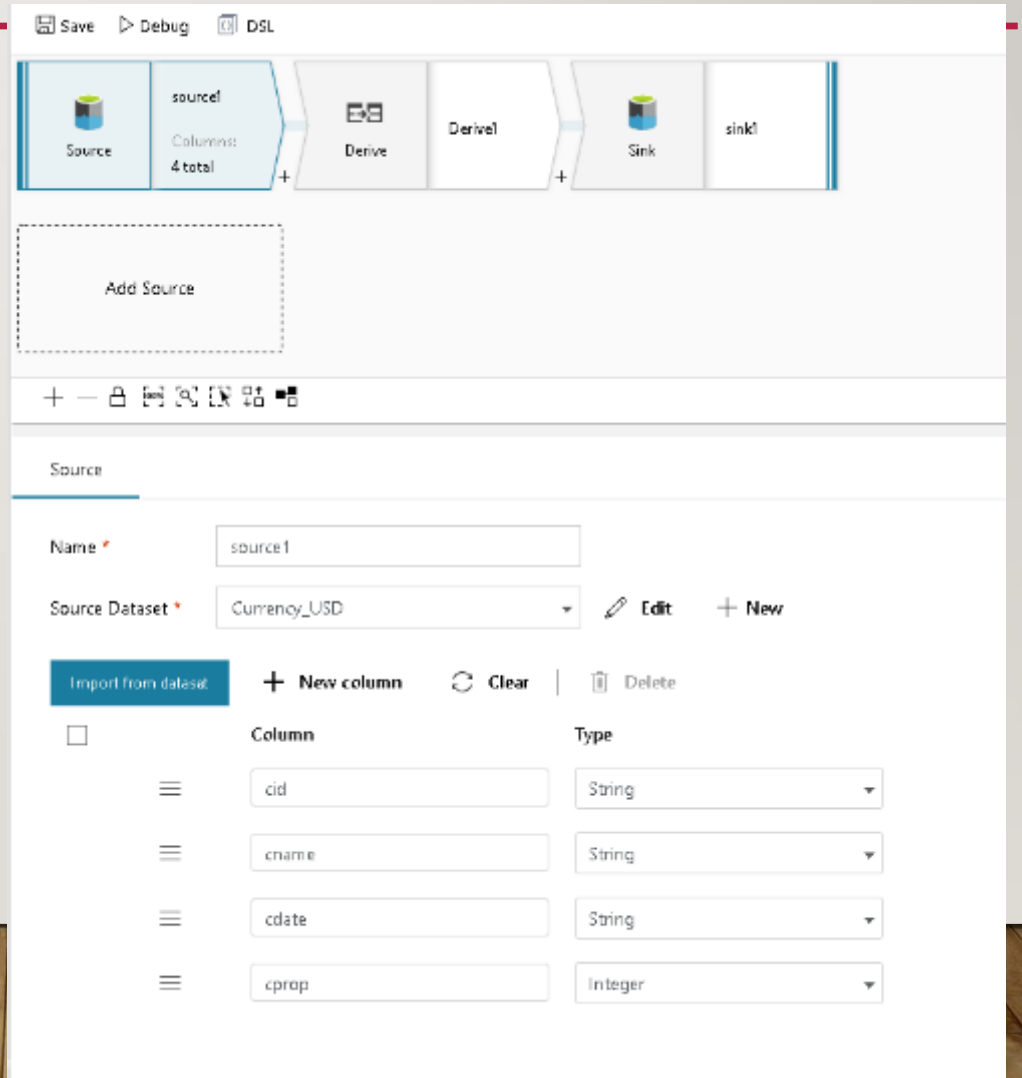

HANDLING SOURCE CHANGES

t: @bizdataviz
In/cseferlis

- Data Engineer derives columns using template expression patterns based on name and type matching. No need to define static field names.

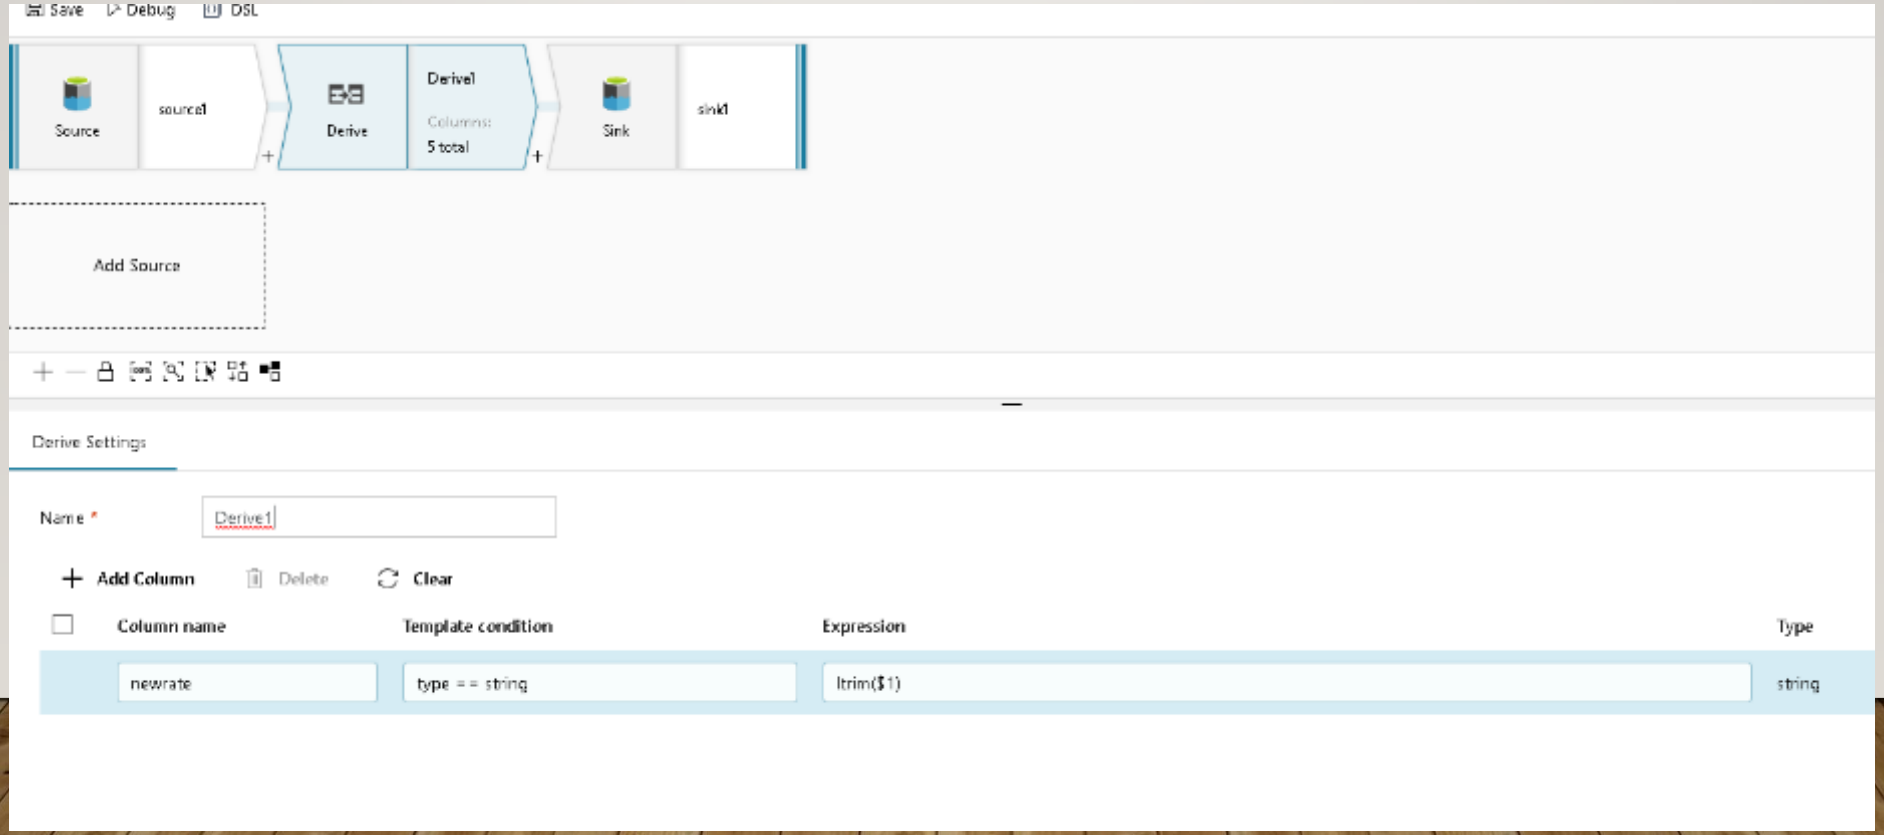

HANDLING SOURCE CHANGES

t: @bizdataviz
ln/cseferlis

- Data Engineer derives columns using template expression based on name and type matching. No need to define static field names.

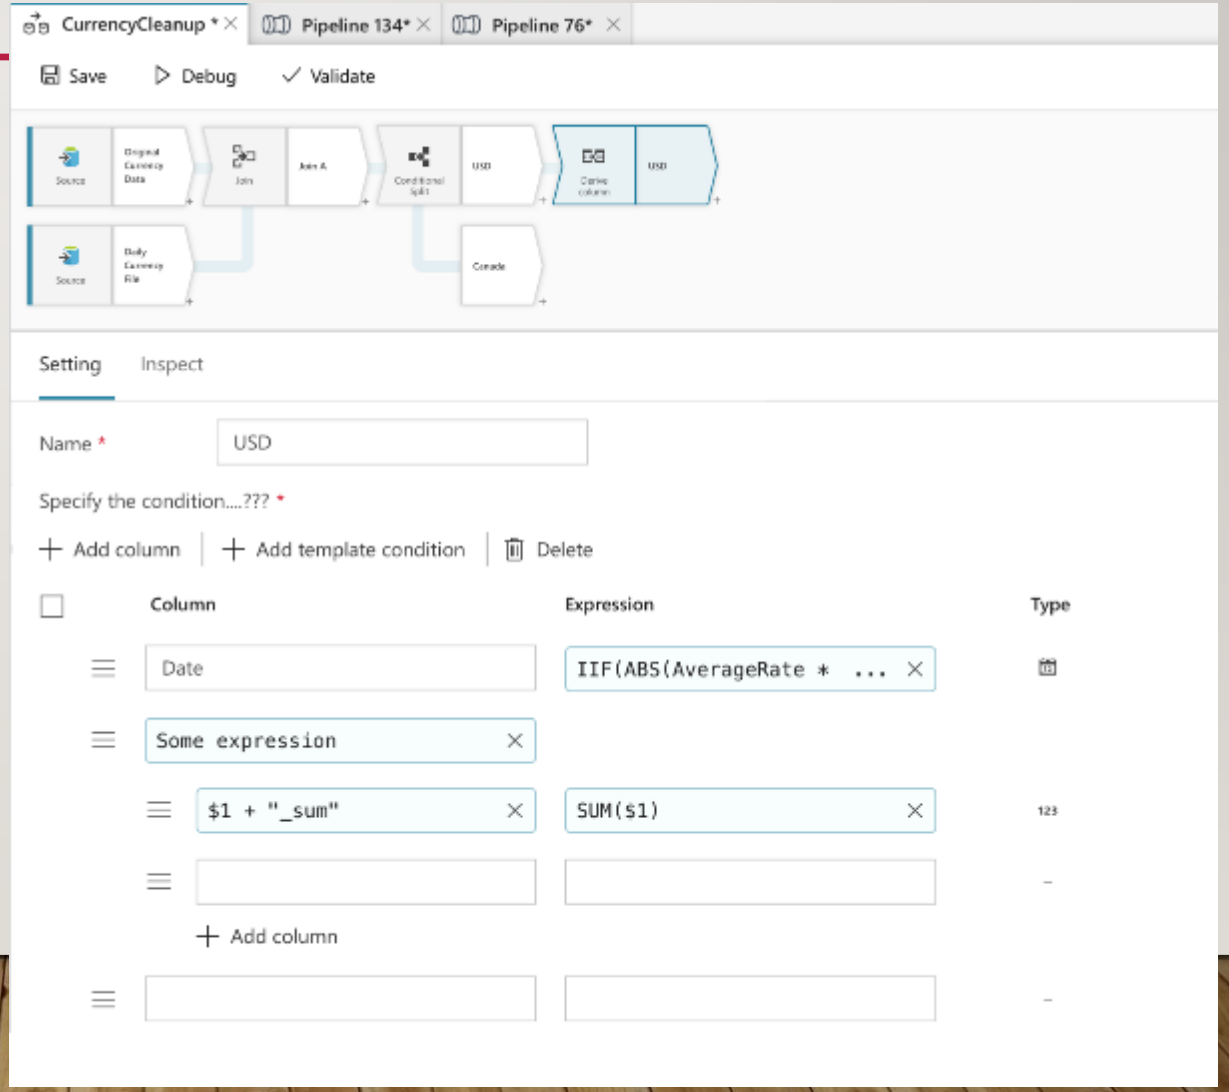

HANDLING SOURCE CHANGES

t: @bizdataviz
In/cseferlis

- Data Engineer derives columns using template expression based on name and type matching

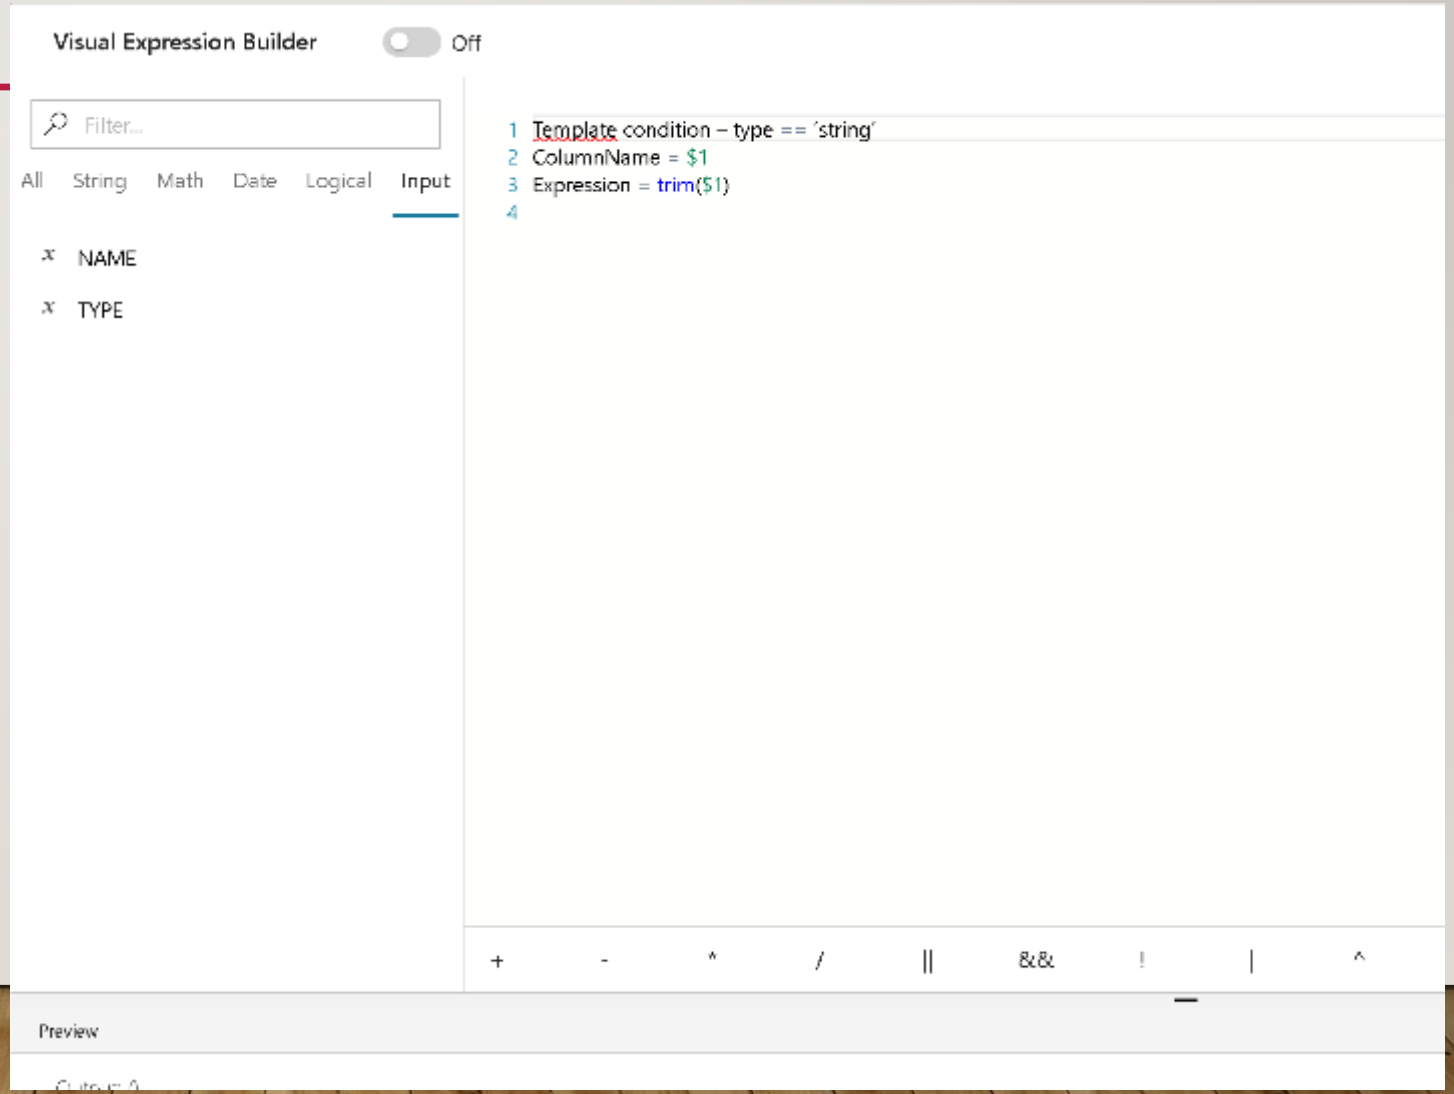

HANDLING SOURCE CHANGES

t: @bizdataviz
ln/cseferlis

- Sink all incoming fields along with new derived field

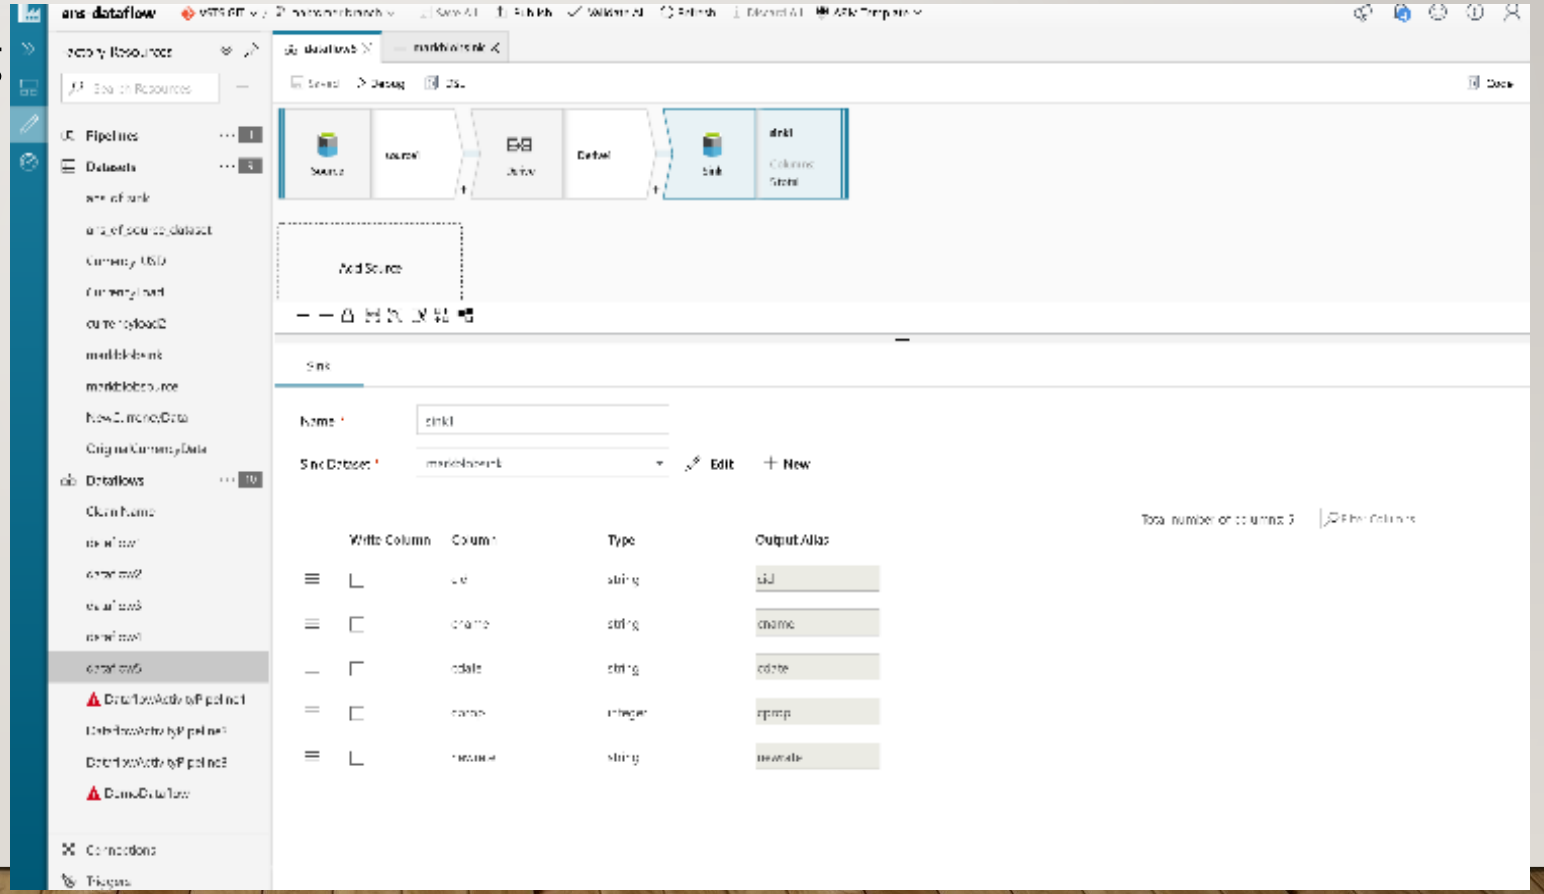

IMPORTANT LINKS:

t: @bizdataviz
In/cseferlis

- Sign-up for ADF Data Flow service:
  - http://aka.ms/dataflowpreview
  - Microsoft Azure must whitelist your subscription ID to turn on the feature for you
- GitHub Repository for documentation:
  - https://github.com/kromerm/adfdataflowdocs/

THANK YOU!