



Azure Stream Analytics



Jack Bender
Industry Architect
US Manufacturing

Unlocking Real-time Insights

Time to Insight is Critical

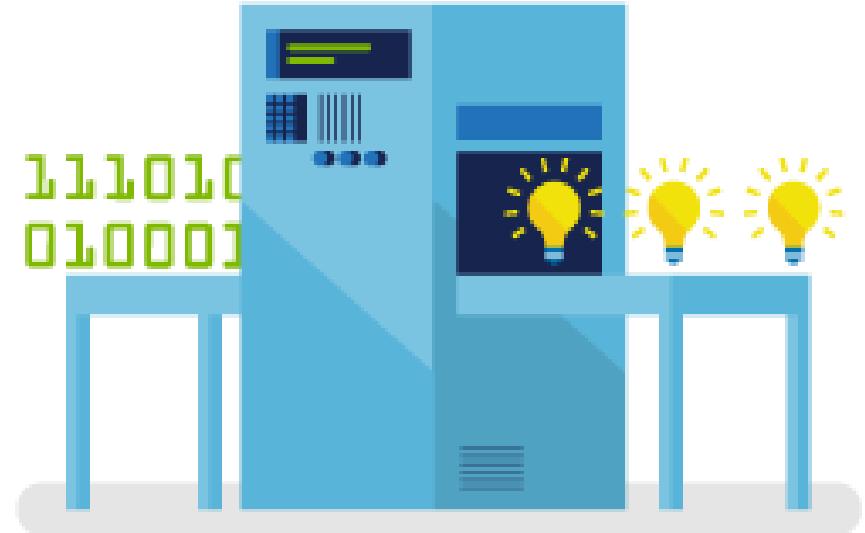
Reducing decision latency can unlock business value

Insights are Perishable

Window of opportunity for insights to be actionable

Ask Questions to Data in Motion

Can't wait for data to get to rest before running computation



Real-time Stream Processing

Simple Event Processing

- Filter
- Transform
- Enrich
- Split
- Route

Event Stream Processing

[Simple event processing] +
Aggregate
Rules

Complex Event Processing

[Event Stream Processing] +
Pattern detection
Time windows
Joins & correlations



Scenario Types

Actions by Human Actors

"See and seize" insights

Live visualization

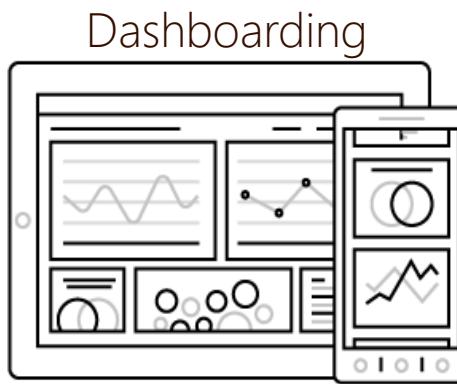
Alerts and alarms

Dynamic aggregation

Machine to Machine Interactions

Data movement with enrichment

Kick-off workflows for automation



Automation



Enriched Data Movement





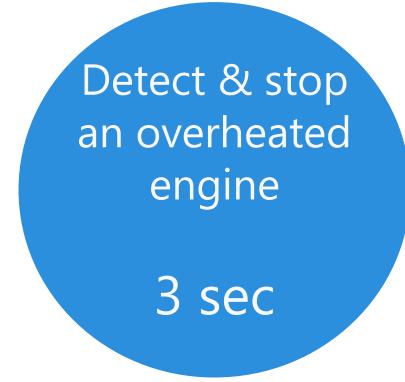
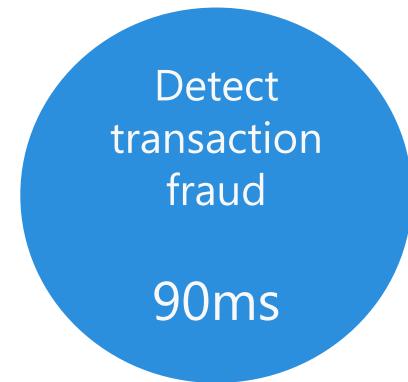
Overview

What is Azure Stream Analytics?

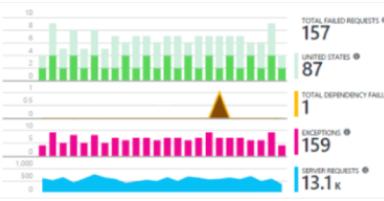
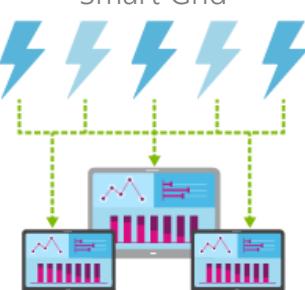
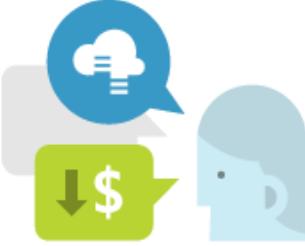
- In-memory Streaming Processing Engine
- Fully Managed: Streaming as a Service!
- Advanced time management made easy



Enables customer to get real-time Insights
when time to action is critical



Stream Analytics is relevant across industries

Real-time Fraud Detection	Streaming ETL	Predictive Maintenance	Call Center Analytics
			
IT Infrastructure and Network Monitoring	Customer Behavior Prediction	Log Analytics	Real-time Cross Sell Offers
			
Fleet monitoring and Connected Cars	Real-time Patient Monitoring	Smart Grid	Real-time Marketing
			

ASA key differentiators



Programmer
productivity: SQL
+ Extensibility



Intelligent Edge,
and Cloud



Serverless and
low TCO



Easy to get
started



Enterprise grade
SLA (99.9%)

Programmer Productivity:

- Simple SQL language
- Extensibility
 - User-Defined Functions (UDF) and Aggregates (UDA)
 - JavaScript (available)
 - C# (preview on edge), more language to come
- Machine-Learning based Anomaly Detection



Programmer
productivity: SQL
+ Extensibility



Intelligent Edge,
and Cloud



Serverless and
low TCO



Easy to get
started



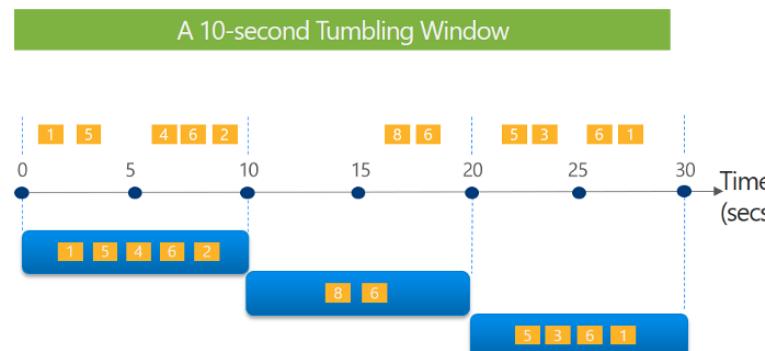
Enterprise grade
SLA (99.9%)

POWERFUL ANALYTICS WITH SIMPLE SQL CONSTRUCTS

- SQL language extended with temporal functions
- Ability to join different streams of data, or streams and reference data
- Time-based windows (between 1 ms and 7 days)
- Using GROUP BY operator for an easy integration with SQL

Example: Tumbling Window

Tell me the count of tweets per time zone every 10 seconds



```
SELECT TimeZone, COUNT(*) AS Count
FROM TwitterStream TIMESTAMP BY CreatedAt
GROUP BY TimeZone, TumblingWindow(second,10)
```

Stream Analytics Query Language (SAQL)

Declarative SQL like language to describe transformations

Filters ("Where")

Projections ("Select")

Time-window and property-based aggregates
("Group By")

Time-shifted joins (specifying time bounds within which the joining events must occur)

and all combinations thereof

Data Manipulation

SELECT
FROM
WHERE
HAVING
GROUP BY
CASE WHEN THEN
ELSE
INNER/LEFT OUTER JOIN
UNION
CROSS/OUTER APPLY
CAST INTO
ORDER BY ASC, DSC

Aggregation

SUM
COUNT
AVG
MIN
MAX
STDEV
STDEVP
VAR
VARP
TopOne

Date and Time

DateName
DatePart Day, Month, Year
Concat
DateDiff
DateTimeFromParts
DateAdd

Temporal
Lag
IsFirst
Last
CollectTop

Windowing Extensions

TumblingWindow
HoppingWindow
SlidingWindow

Scaling Extensions

WITH
PARTITION BY
OVER

String

Len
Concat
CharIndex
Substring
Lower, Upper
PatIndex

Mathematical

ABS
CEILING
EXP
FLOOR
POWER
SIGN
SQUARE
SQRT

Geospatial (preview)

CreatePoint
CreatePolygon
CreateLineString
ST_DISTANCE
ST_WITHIN
ST_OVERLAPS
ST_INTERSECTS

Differentiators

1,915 lines of code with Apache Storm

```
@ApplicationAnnotation(name="WordCountDemo")
public class Application implements StreamingApplication
{
    protected String fileName =
    "com/datatorrent/demos/wordcount/samplefile.txt";
        private Locality locality = null;

    @Override public void populateDAG(DAG dag, Configuration
conf)
    {
        locality = Locality.CONTAINER_LOCAL;
        WordCountInputOperator input =
        dag.addOperator("wordinput", new
        WordCountInputOperator());
        input.setFileName(fileName);
        UniqueCounter<String> wordCount =
        dag.addOperator("count", new
        UniqueCounter<String>());
        dag.addStream("wordinput-count", input.outputPort,
        wordCount.data).setLocality(locality);
        ConsoleOutputOperator consoleOperator =
        dag.addOperator("console", new
        ConsoleOutputOperator());
        dag.addStream("count-console",wordCount.count,
        consoleOperator.input);
    }
}
```

3 lines of SQL in Azure Stream Analytics

```
SELECT Avg(Purchase), ScoreTollId, Count(*)
FROM GameDataStream
GROUP BY TumblingWindows(5, Minute), Score
```

Intelligent Edge and Cloud

- Same queries running either on Edge or Cloud
- Leveraging IoT Edge for security and deployment
- No equivalent service from our main competitors
- Working on Hybrid jobs



Programmer
productivity: SQL
+ Extensibility



Intelligent Edge,
and Cloud



Serverless and
low TCO



Easy to get
started



Enterprise grade
SLA (99.9%)

Serverless and low TCO

- Managed job service: no cluster to manage, just run job
- Pay as you go: pay by hour, start, stop, scale when needed – no commitment
- Scale to 120 SUs and beyond



Programmer productivity: SQL + Extensibility



Intelligent Edge, and Cloud



Serverless and low TCO



Easy to get started



Enterprise grade SLA (99.9%)

Easy to get started

- Native integration with the rest of Azure Ecosystem
- “5 minutes to solution” without writing any code



Programmer productivity: SQL + Extensibility



Intelligent Edge, and Cloud



Serverless and low TCO

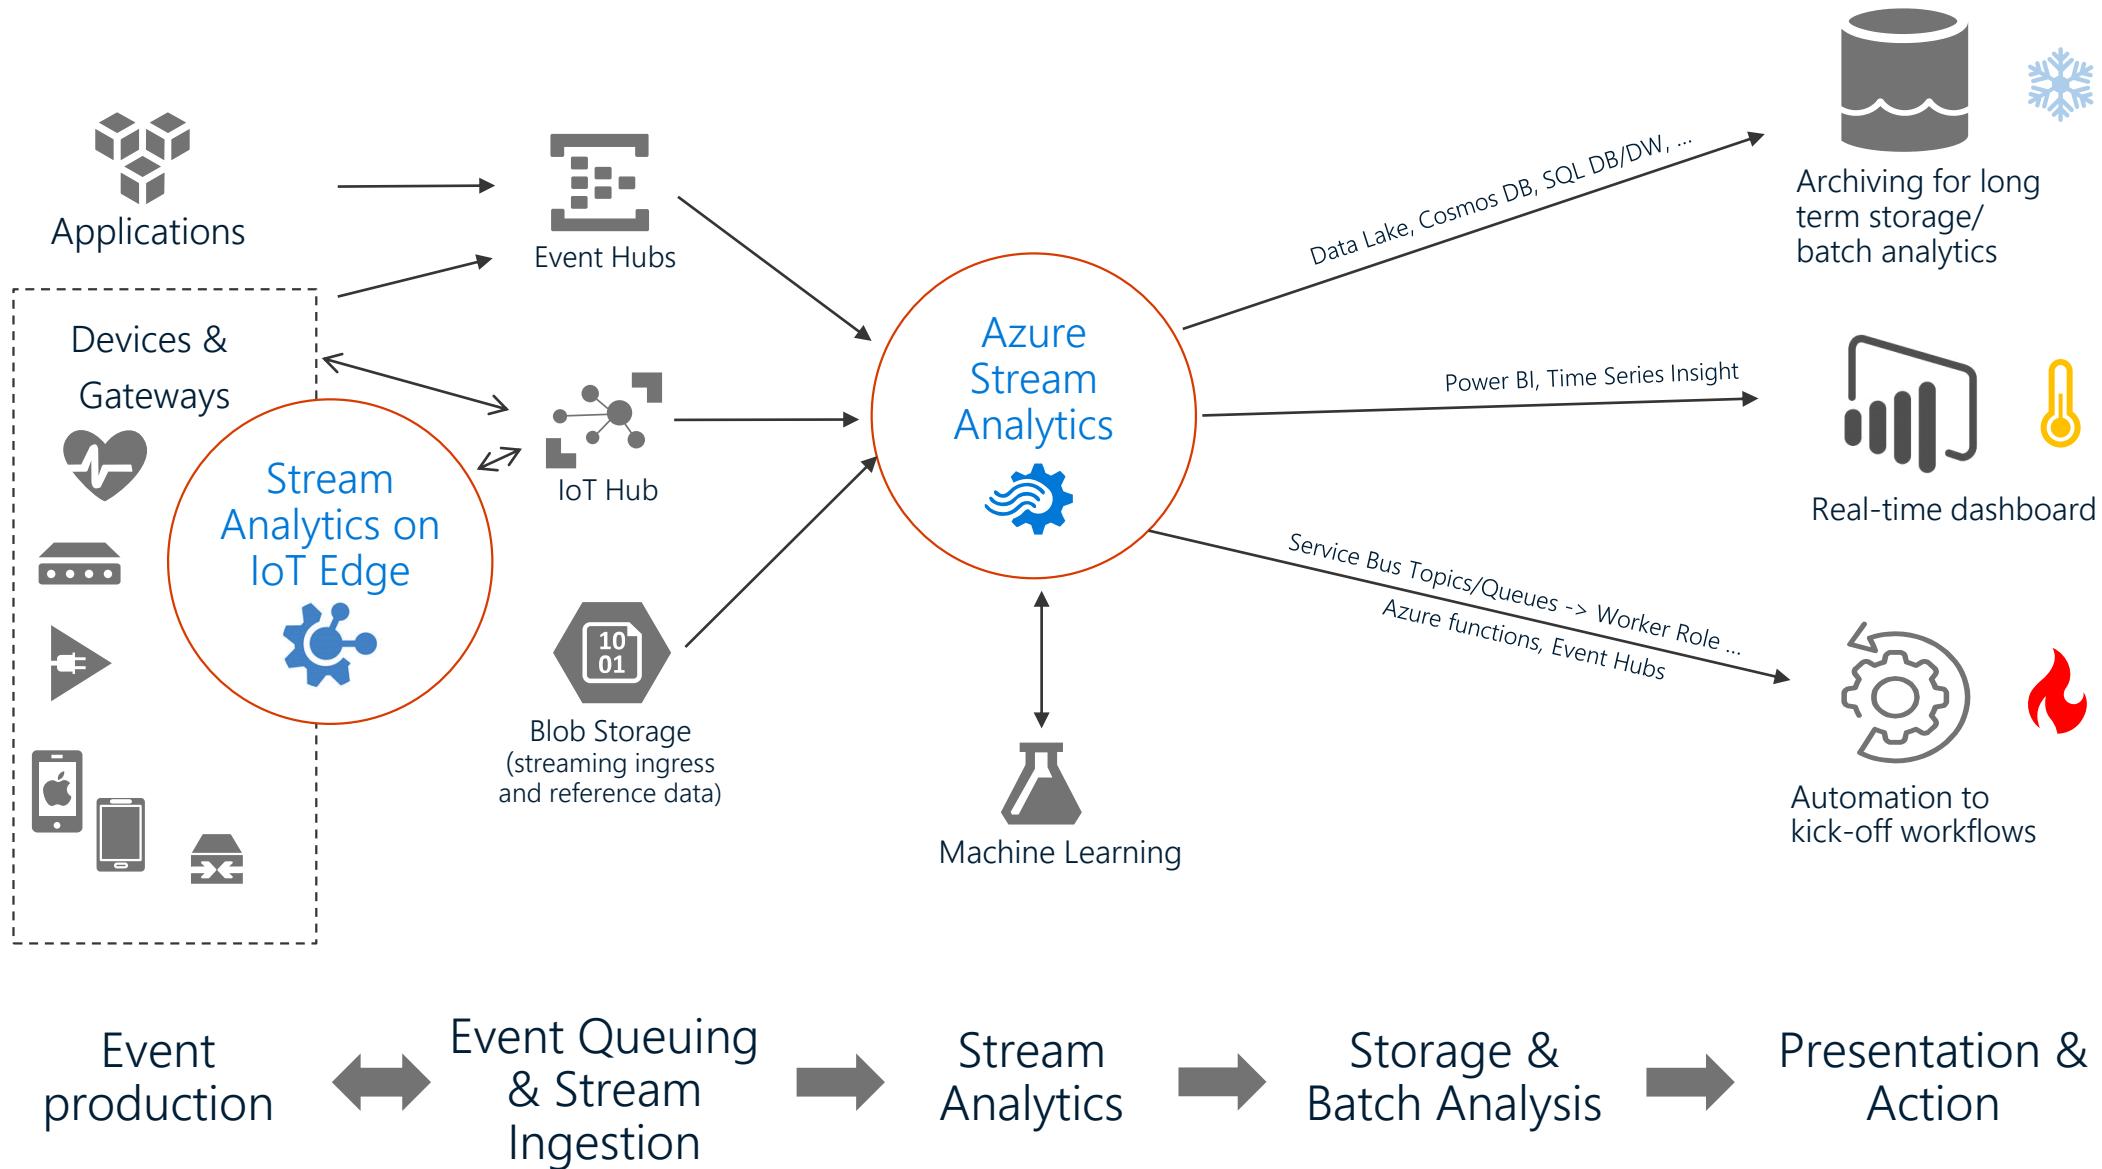


Easy to get started



Enterprise grade SLA (99.9%)

ANALYTICS WITH AZURE STREAM ANALYTICS



Enterprise grade SLA

- Unique SLA in the industry: 99.9% at the minute level
 - Google Dataflow provides 99.5%, AWS Kinesis Analytics doesn't provide SLA
- Working towards 99.99% and helping customers to write pipelines designed to run 24/7



Programmer productivity: SQL + Extensibility



Intelligent Edge, and Cloud



Serverless and low TCO



Easy to get started

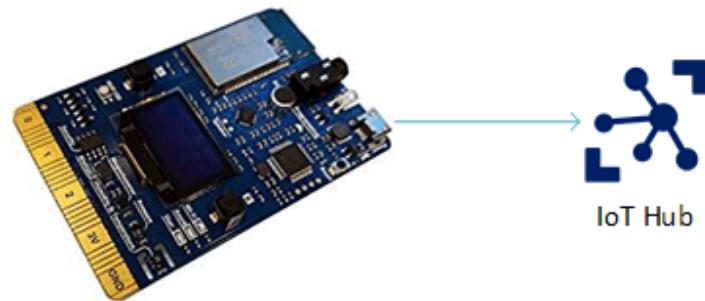


Enterprise grade SLA (99.9%)



Components of ASA

Data Streaming from a Device



```
2/27/2018 5:05:38 PM> Device: [JBAZ3166], Data:  
[{"humidity":37.10,"temp":21.50,"pressure":1014.34,"magnetometerX":994,"magnetometerY":613,"magnetometerZ":1659,  
"accelerometerX":-12,"accelerometerY":13,"accelerometerZ":997,"gyroscopeX":  
-1330,"gyroscopeY":1190,"gyroscopeZ":350}]Properties:  
'timestamp': 'Tue Feb 27 22:05:37 2018'
```

Stream Analytics Job

Job definition includes inputs, a query, and output

Inputs are from where the job reads the data stream

Query runs for perpetuity unless explicitly stopped and transforms the input stream

Output is where the job sends the job results to

Inputs	Query
2	<pre>1 WITH Telemetry AS (2 --MX Chip Telemetry Data 3 SELECT 4 IoTHub.ConnectionDeviceId as Device, 5 System.TimeStamp AS WindowEnd, 6 AVG(Stream.humidity) as humidity, 7 AVG(Stream.temp) as TempC, 8 AVG((Stream.temp*1.8)+32)) as TempF, 9 AVG(Stream.pressure) as pressure, 10 AVG(Stream.magnetometerX) as magnetometerX, 11 AVG(Stream.magnetometerY) as magnetometerY, 12 AVG(Stream.magnetometerZ) as magnetometerZ, 13 MIN(Stream.accelerometerX) as accelerometerX, 14 MIN(Stream.accelerometerY) as accelerometerY,</pre>
DeviceRulesBlob	
IoTStream	
Outputs	See more
4	
AlertsQueue	
AlertsQueueAnomaly	

Stream Analytics Job Inputs & Outputs



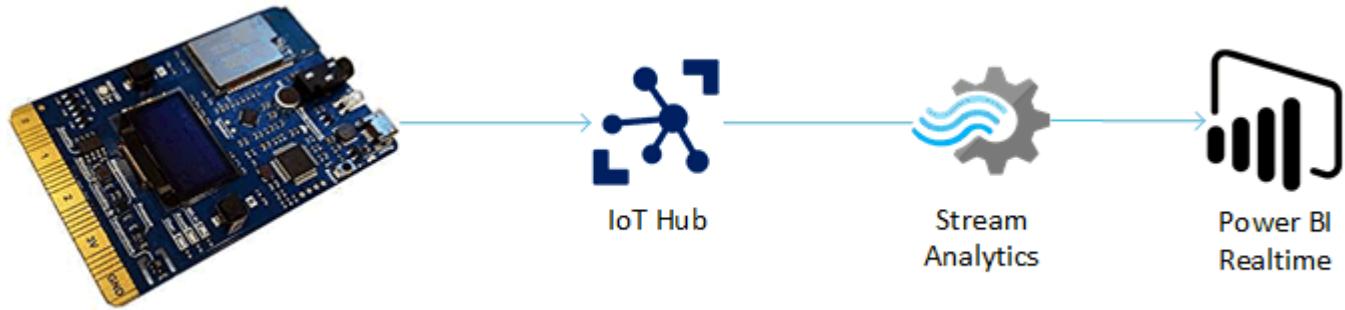
There are two distinct types of Inputs

- Data Streams:
 - IoT Hub
 - Event Hub
 - Azure Blob storage
- Reference data:
 - Azure Blob storage

Data Outputs Supported

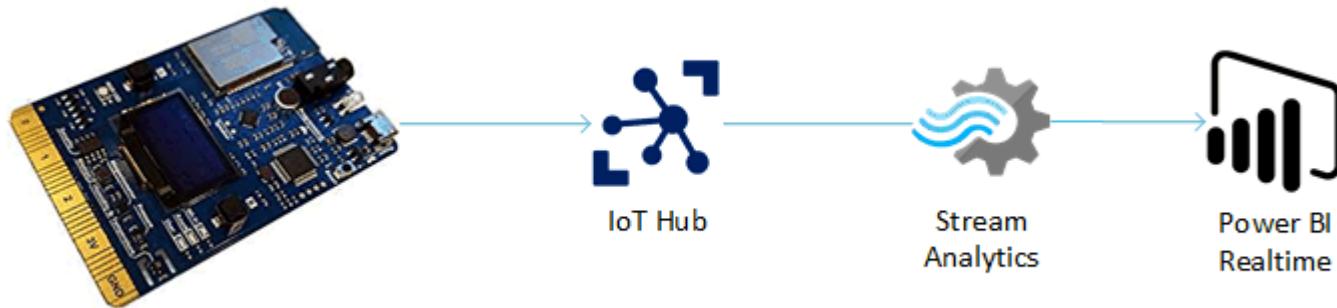
- Azure Data Lake Store
- SQL Database
- Blob storage
- Event Hub
- Power BI
- Table Storage
- Service Bus Queues
- Service Bus Topics
- Azure Cosmos DB
- Azure Functions

Basic Query to Power BI Realtime



```
1 WITH Telemetry AS (
2   --MX Chip Telemetry Data
3   SELECT
4     Stream.*
5   FROM
6     IoTStream Stream
7 )
8
9   SELECT
10  Telemetry.*
11  INTO PowerBI
12  from Telemetry
```

Enhanced Query to Power BI Realtime

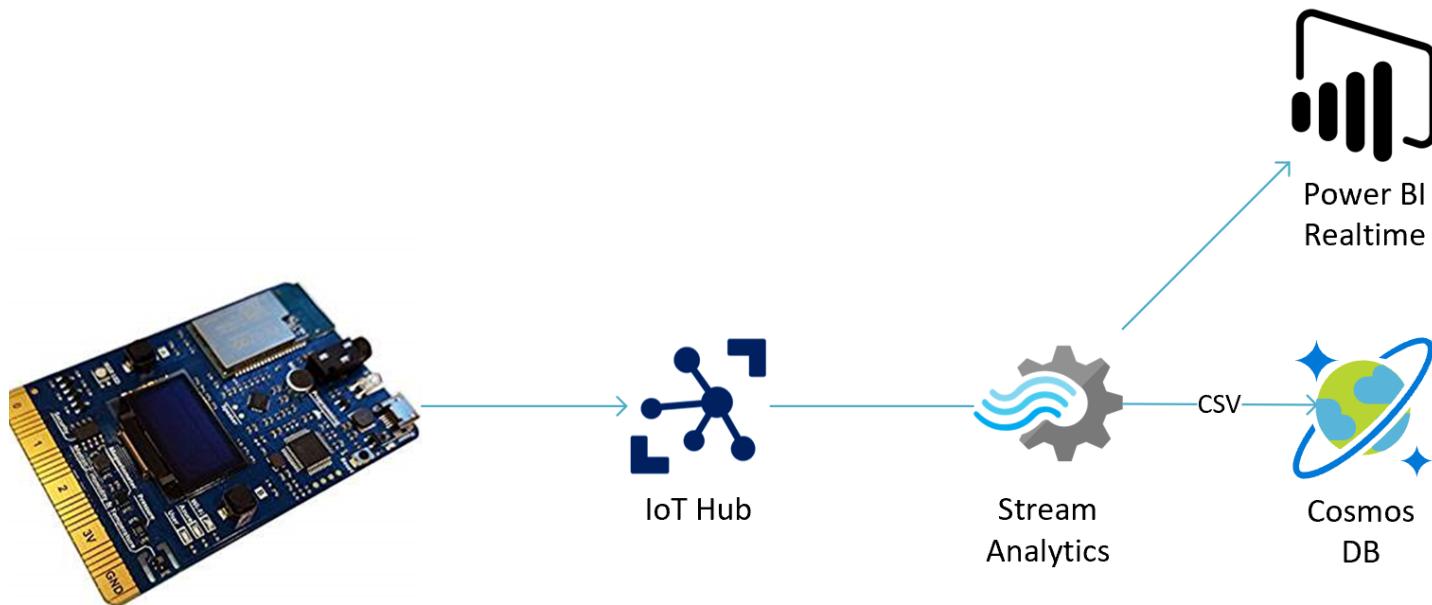


- Device ID/Name
- Time
- Fahrenheit

```
2/27/2018 5:05:38 PM> Device: [JBAZ3166], Data:  
[{"humidity":37.10,"temp":21.50,"pressure":1014.34,"magnetometerX":994,"magnetometerY":13,"magnetometerZ":997,"accelerometerX":-12,"accelerometerY":13,"accelerometerZ":997,"gyroscopeX":1330,"gyroscopeY":1190,"gyroscopeZ":350}]Properties:  
'timestamp': 'Tue Feb 27 22:05:37 2018'
```

```
1 WITH Telemetry AS (  
2   --MX Chip Telemetry Data  
3     SELECT  
4       IoTHub.ConnectionDeviceId as Device,  
5       IoTHub.EnqueueTime as Time,  
6       Stream.humidity,  
7       Stream.temp as TempC,  
8       ((Stream.temp*1.8)+32) as TempF,  
9       Stream.pressure,  
10      Stream.magnetometerX,  
11      Stream.magnetometerY,  
12      Stream.magnetometerZ,  
13      Stream.accelerometerX,  
14      Stream.accelerometerY,  
15      Stream.accelerometerZ,  
16      Stream.gyroscopeX,  
17      Stream.gyroscopeY,  
18      Stream.gyroscopeZ  
19  
20   FROM  
21     IoTStream Stream  
22   )  
23  
24   SELECT  
25     Telemetry.*  
26   INTO PowerBI  
27   from Telemetry
```

Data to Cosmos DB & Power BI Realtime

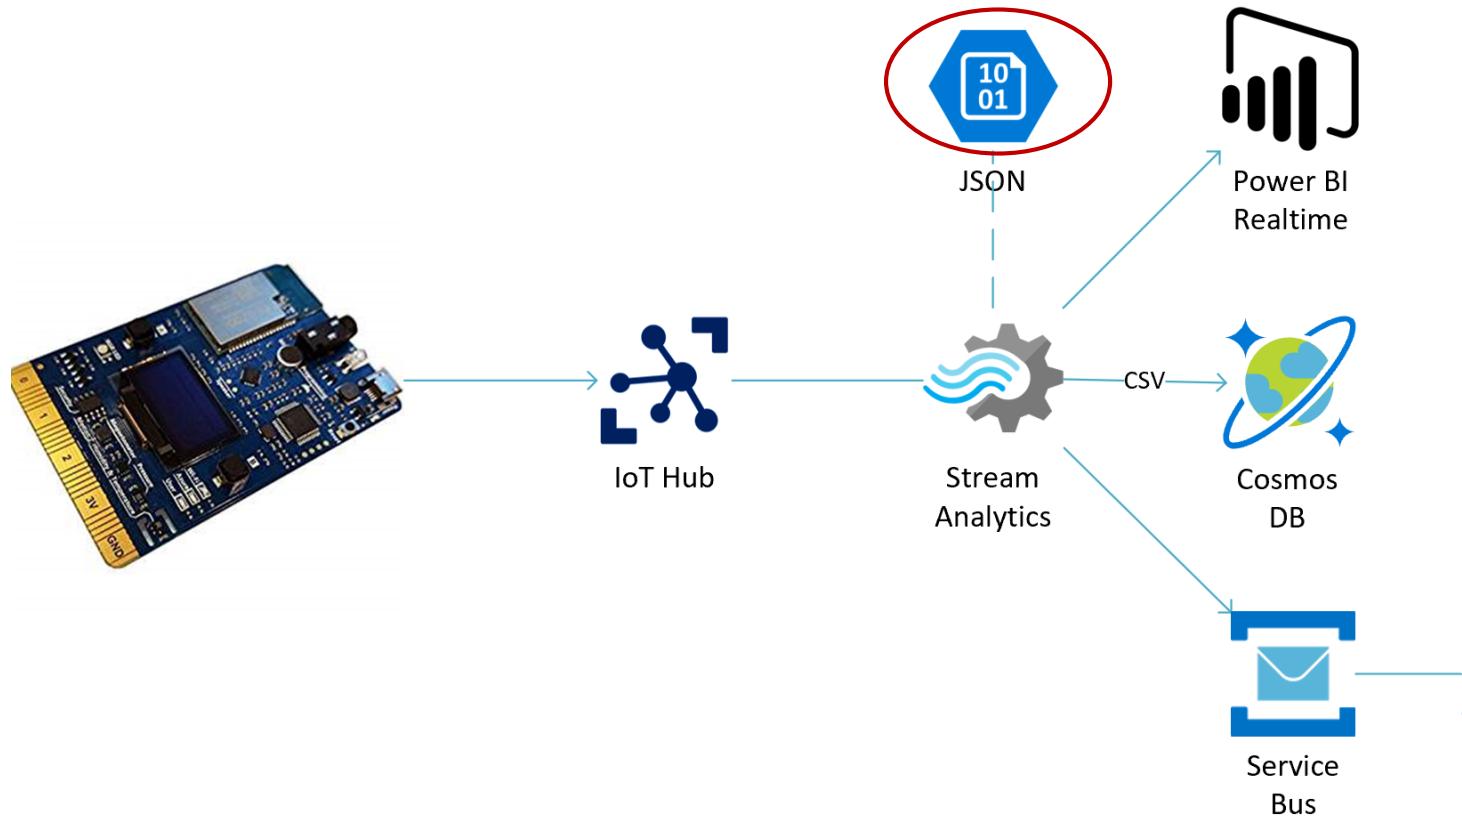


```
24  -- Send all data to Cosmos DB
25  SELECT
26  Telemetry.*
27  INTO CosmosDB
28  from Telemetry
29
30  -- Send some data to Power BI
31  SELECT
32  Telemetry.Device,
33  Telemetry.Time,
34  Telemetry.accelerometerX as 'Accelerometer X',
35  Telemetry.accelerometerY as 'Accelerometer Y',
36  Telemetry.accelerometerZ as 'Accelerometer Z',
37  Telemetry.TempF as Fahrenheit,
38  Telemetry.TempC as Celsius
39
40  INTO PowerBI
41  from Telemetry
```

Cosmos DB, Power BI Realtime and Text Alert

{ } devicerules.json x

```
1 [{"DeviceType": "Accelerometer", "AccelerometerX": 0, "AccelerometerZ": 0, "AccelerometerXRuleOutput": "XAlert", "AccelerometerZRuleOutput": "ZAlert"}]
```



```
1 27  SELECT
2 28  Telemetry.*
3 29  INTO CosmosDB
4 30  from Telemetry
5
6 32  -- Send some data to Power BI
7 33  SELECT
8 34  Telemetry.Device,
9 35  Telemetry.WindowEnd as Time,
10 36  Telemetry.accelerometerX as 'Accelerometer X',
11 37  Telemetry.accelerometerY as 'Accelerometer Y',
12 38  Telemetry.accelerometerZ as 'Accelerometer Z',
13 39  Telemetry.TempF as Fahrenheit,
14 40  Telemetry.TempC as Celsius
15
16 42  INTO PowerBI
17 43  from Telemetry
18
19 45  -- Send Alerts to Message Queue
20 46  SELECT
21 47  Telemetry.Device,
22 48  Telemetry.WindowEnd as Time,
23 49  Telemetry.accelerometerZ as 'Accelerometer Z',
24 50  Ref.AccelerometerZ as 'Z Threshold',
      Ref.AccelerometerZRuleOutput
51 52  INTO AlertsQueue
53
54 53  from Telemetry
55 54  JOIN DeviceRulesBlob Ref ON Ref.DeviceType = 'Accelerometer'
56 55  WHERE Telemetry.accelerometerZ < Ref.AccelerometerZ
```



Time Functions

Events and Time

Every event that flows through the system has a timestamp

ASA supports:

Arrival Time - Event timestamps based on arrival time (input adapter clock, e.g., Event Hubs)

App Time - Event timestamps based on a timestamp field in the actual event tuple

User can pick up App Time from the payload

SELECT * FROM EntryStream TIMESTAMP BY EntryTime

System can assign timestamps automatically based on the event arrival time

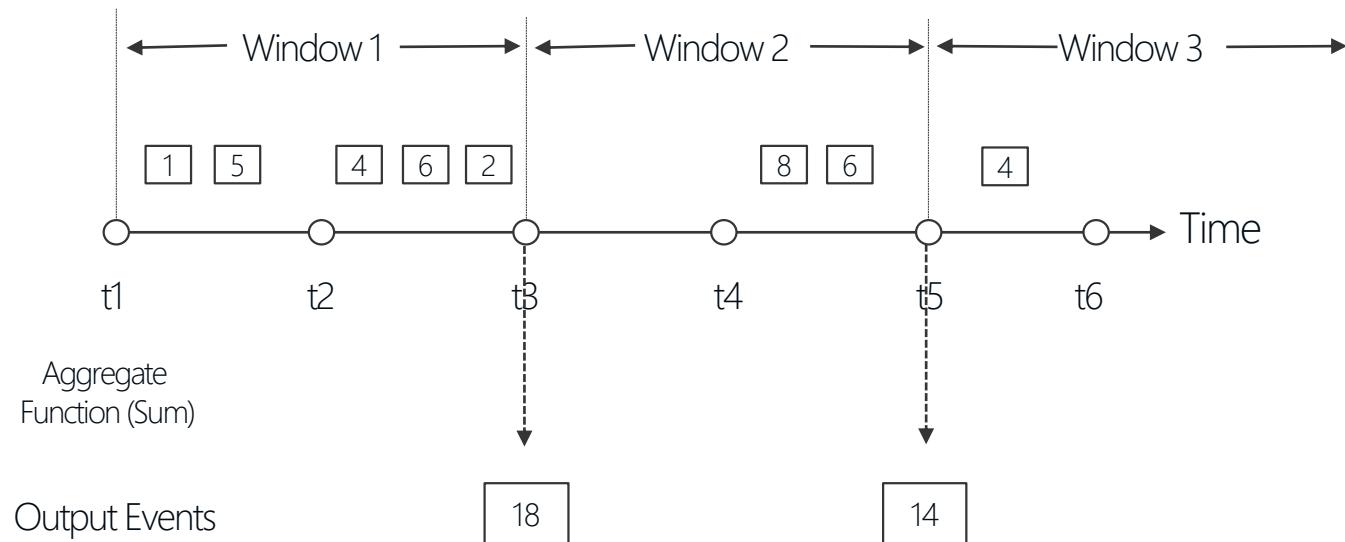
SELECT * FROM EntryStream

Windowing Concepts

Output at the end of each window

Windows are fixed length

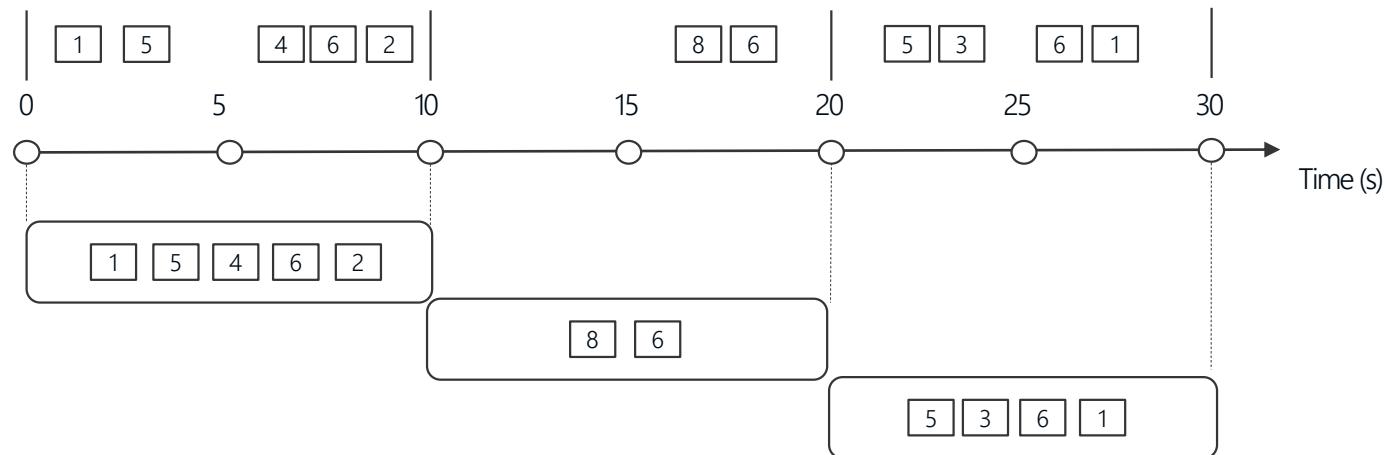
Used in a GROUP BY clause



Tumbling Windows

Every 10 seconds give me the count of vehicles entering each toll booth over the last 10 seconds

A 10-second Tumbling Window

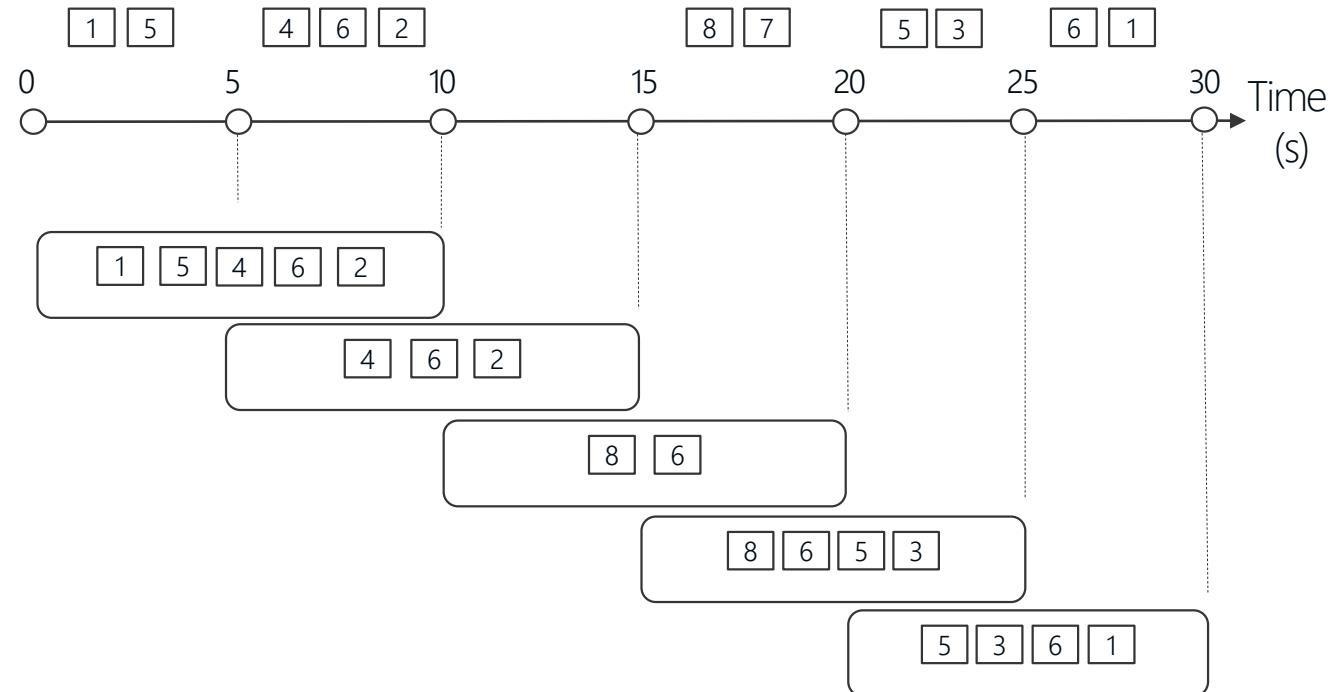


```
SELECT TollId, Count(*)  
FROM EntryStream TIMESTAMP BY EntryTime  
GROUP BY TollId, TumblingWindow(second, 10)
```

Hopping Windows

Every 5 seconds give me the count of vehicles entering each toll booth over the last 10 seconds

A 10 second Hopping Window with a 5 second hop



```
SELECT TollId, Count(*)
```

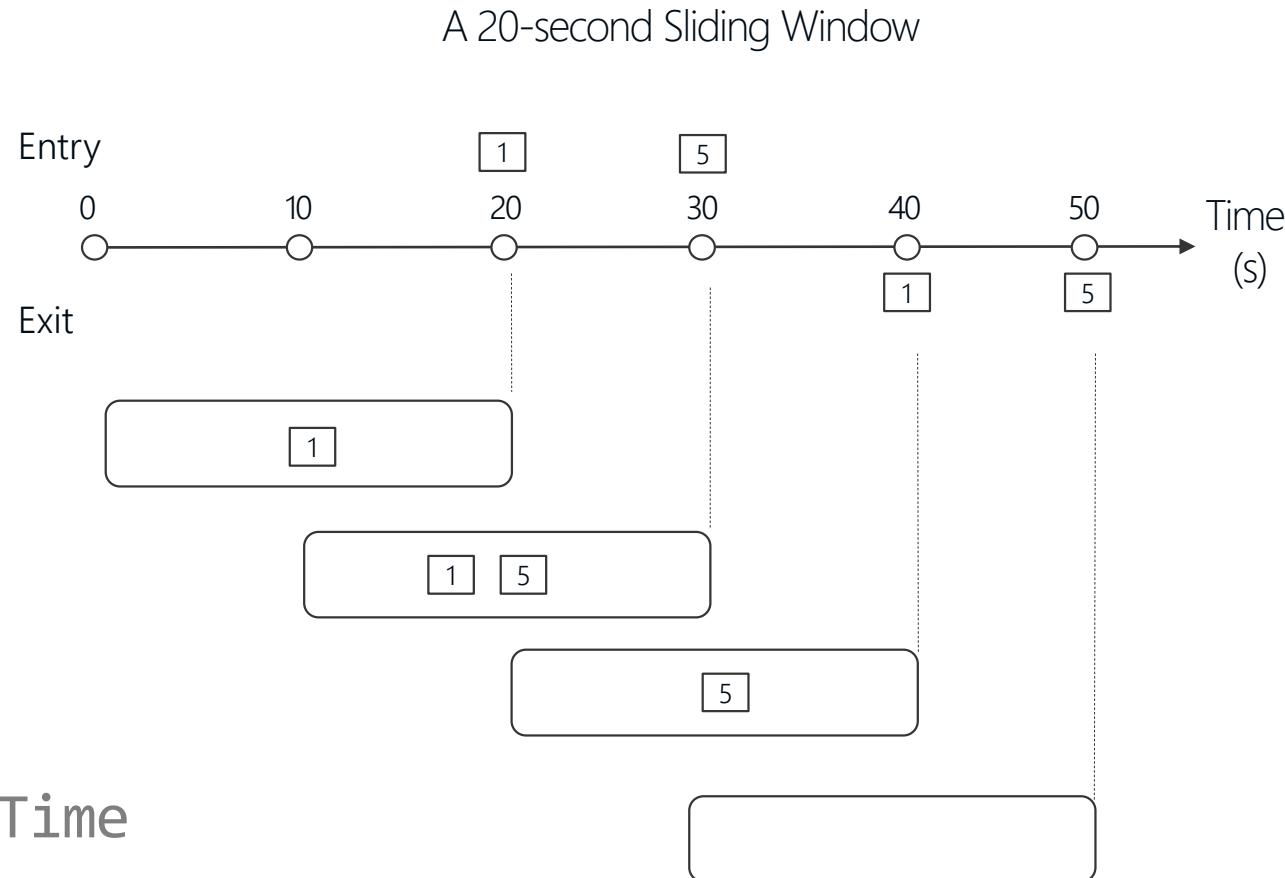
```
FROM EntryStream TIMESTAMP BY EntryTime
```

```
GROUP BY TollId, HoppingWindow(second, 10, 5)
```

Sliding Windows

Find all toll booths that have
served more than 10 vehicles in
the last 20 seconds

```
SELECT TollId, Count(*)  
FROM EntryStream TIMESTAMP BY EntryTime  
GROUP BY TollId, SlidingWindow(second, 20)  
HAVING Count(*) > 10
```



An output is generated whenever an event either enters/leaves the system



Machine Learning Integration

Azure Machine Learning Callouts

Perform real-time scoring on streaming data

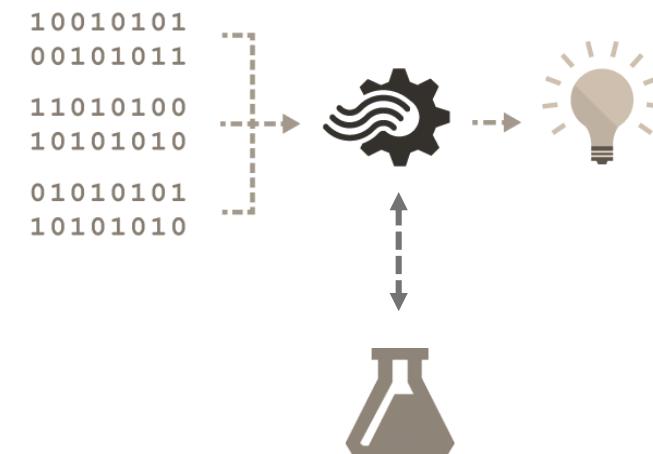
Anomaly Detection and Sentiment Analysis are common use cases

Function calls from the query

Azure ML can publish web endpoints for operationalized ML models

Azure Stream Analytics binds custom function names to such web endpoints

```
SELECT text, sentiment(text) AS score  
FROM myStream
```



Azure Machine Learning Callouts

```
1 WITH AlertData AS (
2   -- Web Simulator Devices
3   SELECT
4     Stream.EventToken AS ID,
5     Stream.DeviceID,
6     Stream.Temperature AS reading,
7     Stream.EventEnqueuedUtcTime AS [time]
8   FROM IoTStream Stream
9
10  WHERE
11    (Stream.EventToken IS NOT NULL) and (Stream.DeviceID > 0)
12
13 UNION
14
15  -- MX Chip
16  SELECT
17    Stream.messageId AS ID,
18    Stream.DeviceID,
19    Stream.Temperature AS reading,
20    Stream.EventEnqueuedUtcTime AS [time]
21  FROM IoTStream Stream
22
23 WHERE
24    (Stream.messageId IS NOT NULL) and (Stream.DeviceID > 0)
25
26 UNION
27
28  -- Particle Device
29  SELECT
30    Stream.published_at AS ID,
31    Stream.device_id AS DeviceID,
32    CAST(Stream.data AS float) AS reading,
33    Stream.EventEnqueuedUtcTime AS [time]
34
35 FROM IoTStream Stream
36
37 WHERE
38   CAST(Stream.data AS float) > 0 and Stream.Event='temperature'
39
40 ),
```

```
42 AmlStep AS (
43   SELECT
44     ID,
45     DeviceID,
46     reading,
47     time,
48     aml(ID, DeviceID, reading, time) AS result
49   FROM AlertData
50
51 )
52
53 SELECT
54   result.ID,
55   result.DeviceID,
56   'Temperature' AS ReadingType,
57   result.reading,
58   CAST(result.time AS datetime) AS time,
59   result.[abn]
60
61 --INTO testoutput
62 INTO AlertsQueue
63 FROM AmlStep
64 WHERE CAST(result.[abn] AS bigint)=1
65 AND (LAG(result.DeviceID) OVER (PARTITION BY result.DeviceID, result.reading LIMIT DURATION(minute, 1)) IS NULL)
```



Demo

SpikeAndDipScores in Azure Stream Analytics

```
1 WITH Telemetry AS (
2   --MX Chip Telemetry Data
3   SELECT
4     IoTHub.ConnectionDeviceId as Device,
5     System.TimeStamp AS WindowEnd,
6     AVG(Stream.humidity) as humidity,
7     AVG(Stream.temp) as TempC,
8     AVG((Stream.temp*1.8)+32)) as TempF,
9     AVG(Stream.pressure) as pressure,
10    AVG(Stream.magnetometerX) as magnetometerX,
11    AVG(Stream.magnetometerY) as magnetometerY,
12    AVG(Stream.magnetometerZ) as magnetometerZ,
13    MIN(Stream.accelerometerX) as accelerometerX,
14    MIN(Stream.accelerometerY) as accelerometerY,
15    MIN(Stream.accelerometerZ) as accelerometerZ,
16    AVG(Stream.gyroscopeX) as gyroscopeX,
17    AVG(Stream.gyroscopeY) as gyroscopeY,
18    AVG(Stream.gyroscopeZ) as gyroscopeZ
19
20  FROM
21    IoTStream Stream TIMESTAMP BY IoTHub.EnqueueTime
22  GROUP BY IoTHub.ConnectionDeviceId, TumblingWindow(second, 15)
23
24 ),
25
26  -- Anomaly Detection Spike-Dip
27  AnomalyDetectionStep AS (
28  SELECT
29    *,
30    system.timestamp as anomalyDetectionStepTimestamp,
31    AnomalyDetection_SpikeAndDip(accelerometerZ, 99, 20, 'spikesanddips') OVER(LIMIT DURATION(mi, 10)) as SpikeAndDipScores
32  FROM Telemetry
33 )
```

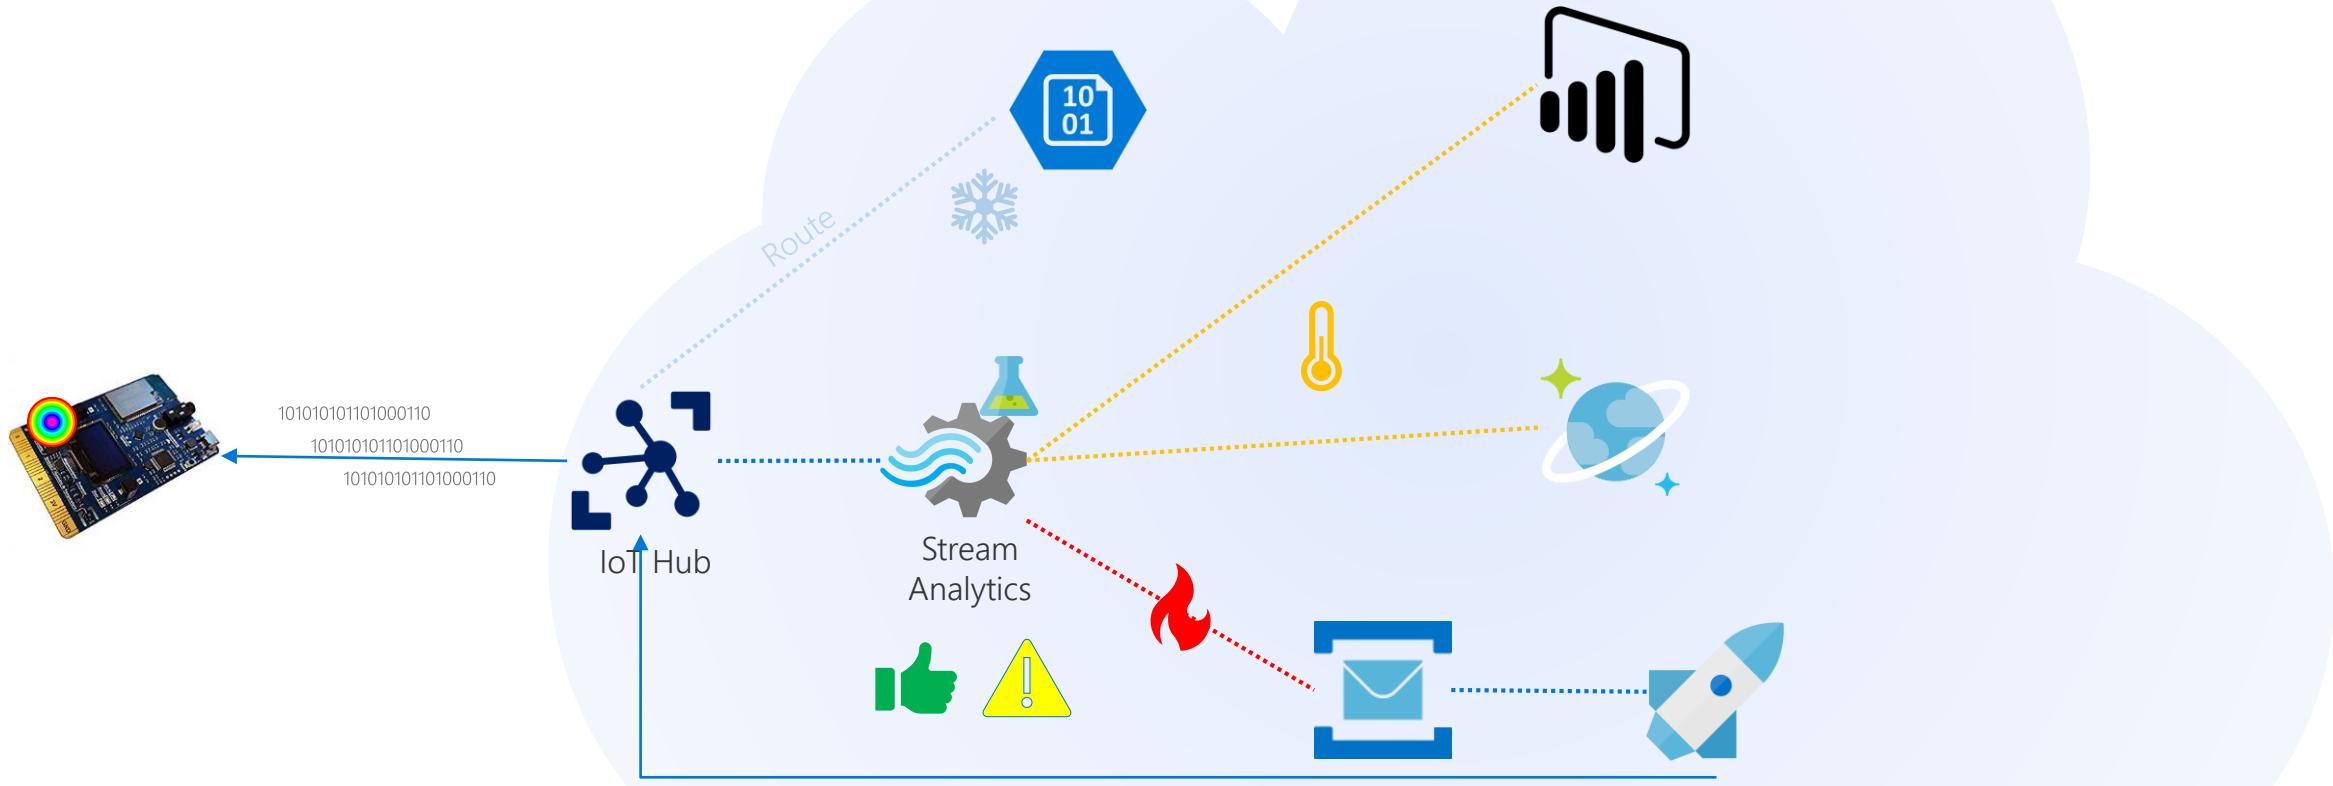
It is recommended to have at least 50 events in each window for best results.

SpikeAndDipScores in Azure Stream Analytics

```
58 -- Send some data to Power BI
59 SELECT
60     Device,
61     WindowEnd as Time,
62     accelerometerX as 'Accelerometer X',
63     accelerometerY as 'Accelerometer Y',
64     accelerometerZ as 'Accelerometer Z',
65     TempF as Fahrenheit,
66     TempC as Celsius,
67     CAST(GetRecordPropertyValue(SpikeAndDipScores, 'Score') as float) as SpikeAndDipScore,
68     CAST(GetRecordPropertyValue(SpikeAndDipScores, 'IsAnomaly') as bigint) as IsSpikeAndDipAnomaly
69
70 INTO PowerBI
71 from AnomalyDetectionStep
72
73 -- Send Alerts to Message Queue
74
75 SELECT
76     Device,
77     WindowEnd as Time,
78     'Accelerometer Z' AS ReadingType,
79     accelerometerZ as reading,
80     CAST(GetRecordPropertyValue(SpikeAndDipScores, 'IsAnomaly') as bigint) as abn
81
82 INTO AlertsQueueAnomaly
83 FROM AnomalyDetectionStep
84 WHERE CAST(GetRecordPropertyValue(SpikeAndDipScores, 'IsAnomaly') as float) = 1
```

It is recommended to have at least 50 events in each window for best results.

SpikeAndDip Architecture





JavaScript user-defined functions

2

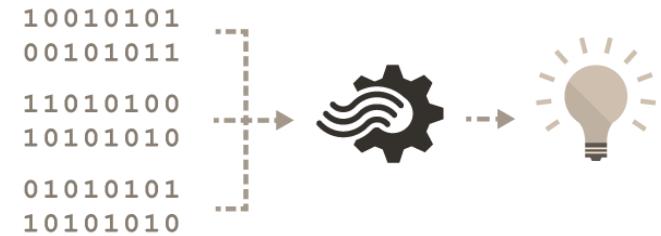
JavaScript user-defined functions

Scenarios where you might find JavaScript user-defined functions useful:

- Parsing and manipulating strings that have regular expression functions, for example, `Regexp_Replace()` and `Regexp_Extract()`
- Decoding and encoding data, for example, binary-to-hex conversion
- Performing mathematic computations with JavaScript `Math` functions
- Performing array operations like sort, join, find, and fill

Here are some things that you cannot do with a JavaScript user-defined function in Stream Analytics:

- Call out external REST endpoints, for example, performing reverse IP lookup or pulling reference data from an external source
- Perform custom event format serialization or deserialization on inputs/outputs
- Create custom aggregates



JavaScript user-defined functions example

```
{  
  "devicename": "bigmachine1",  
  "datetime": "2018-03-01T08:09:09.0470575-05:00",  
  "parameter": "pressure",  
  "value": 12  
}  
  
{  
  "devicename": "bigmachine1",  
  "datetime": "2018-03-01T08:09:09.0470575-05:00",  
  "parameter": "gigathings",  
  "value": 4  
}  
  
{  
  "devicename": "bigmachine1",  
  "datetime": "2018-03-01T08:09:09.0470575-05:00",  
  "parameter": "quarks",  
  "value": 95  
}  
  
{  
  "devicename": "bigmachine1",  
  "datetime": "2018-03-01T08:09:09.0470575-05:00",  
  "parameter": "doodads",  
  "value": 12  
}
```

The JavaScript UDF function takes an array of json objects and flattens it:

```
function main(arr) {  
  //initialize the result object as a blank object  
  var result = {};  
  if (Object.prototype.toString.call(arr) == "[object Array]") {  
    //iterate through all the objects in the array:  
    for (i = 0; i < arr.length; i++) {  
      //get the value of the datetime parameter and add it to the result array  
      //if we haven't already  
      if (result["datetime"] == null) {  
        result["datetime"] = arr[i].datetime;  
      }  
  
      //get the value of the devicename and add it to the result array  
      //if we haven't already  
      if (result["devicename"] == null) {  
        result["devicename"] = arr[i].devicename;  
      }  
  
      //check to see if we have added a parameter/value pair to the results array  
      //and add it to the results array if it isn't in there  
      if(result[arr[i].parameter]==null) {  
        result[arr[i].parameter]=arr[i].value;  
      }  
    }  
  }  
  //pass the array back to the query  
  return result;  
}
```



JavaScript user-defined functions example

The ASA query

```
WITH EventCollection AS
(
  SELECT collect() AS allEvents
  FROM iotevents TIMESTAMP BY datetime
  GROUP BY System.Timestamp, devicename
),
PivotRecord AS
(
  SELECT udf.getValues(allEvents) AS record
  FROM EventCollection
)
SELECT record.*
INTO iotoutput
FROM PivotRecord
```

The resulting file is a CSV blob:

```
datetime,devicename,pressure,quarks,gigathings,doodads
2018-03-01T05:57:42.5690687Z,bigmachine1,11,81,7,50
2018-03-01T05:58:42.7074603Z,bigmachine1,5,61,9,27
2018-03-01T05:59:42.8404016Z,bigmachine1,3,32,9,9
2018-03-01T06:00:42.9714928Z,bigmachine1,11,58,5,3
2018-03-01T06:01:43.1051537Z,bigmachine1,12,30,3,9
2018-03-01T06:02:43.2765973Z,bigmachine1,8,86,12,27
2018-03-01T06:03:43.4033917Z,bigmachine1,6,83,19,41
```





Edge Analytics

Edge Analytics

Local Execution

Stream analytics runs on 'edge devices'

Unlock the Value of Untapped data

Only ~5% of data in industrial processes is sent to the cloud

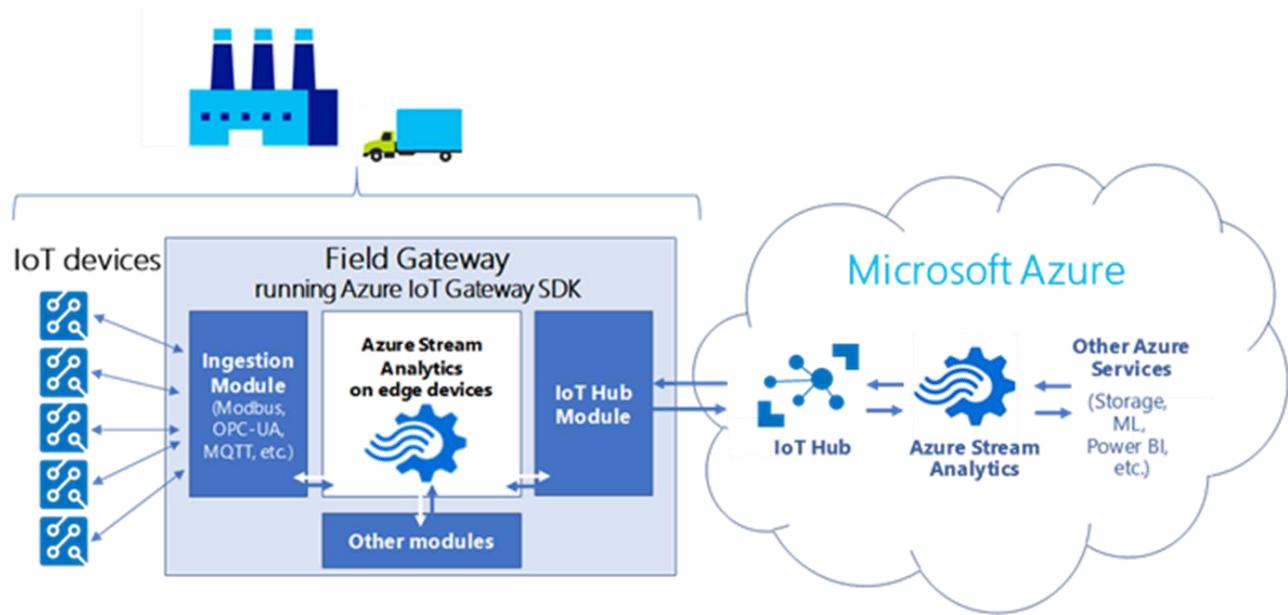
Deploy intelligence near the data to unlock the full value of data

Seamless development and operations

Stream analytics jobs run in the cloud and on edge devices

Intelligent actions

Deploy situational awareness, custom code, ML models on the edge



Edge Analytics Scenarios

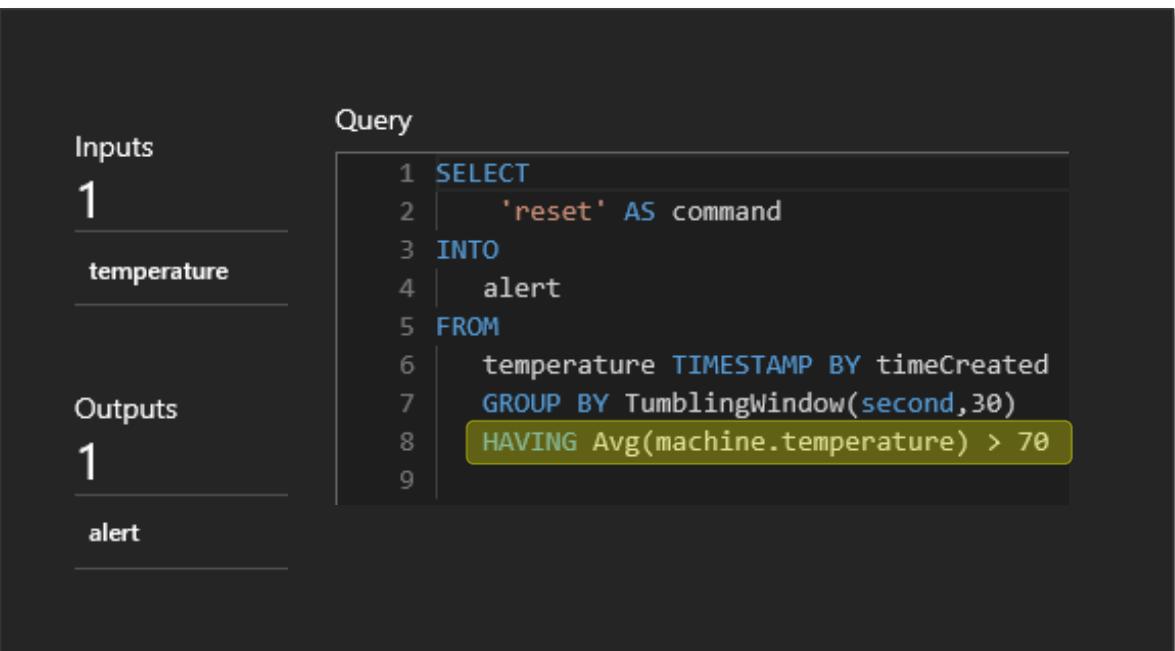
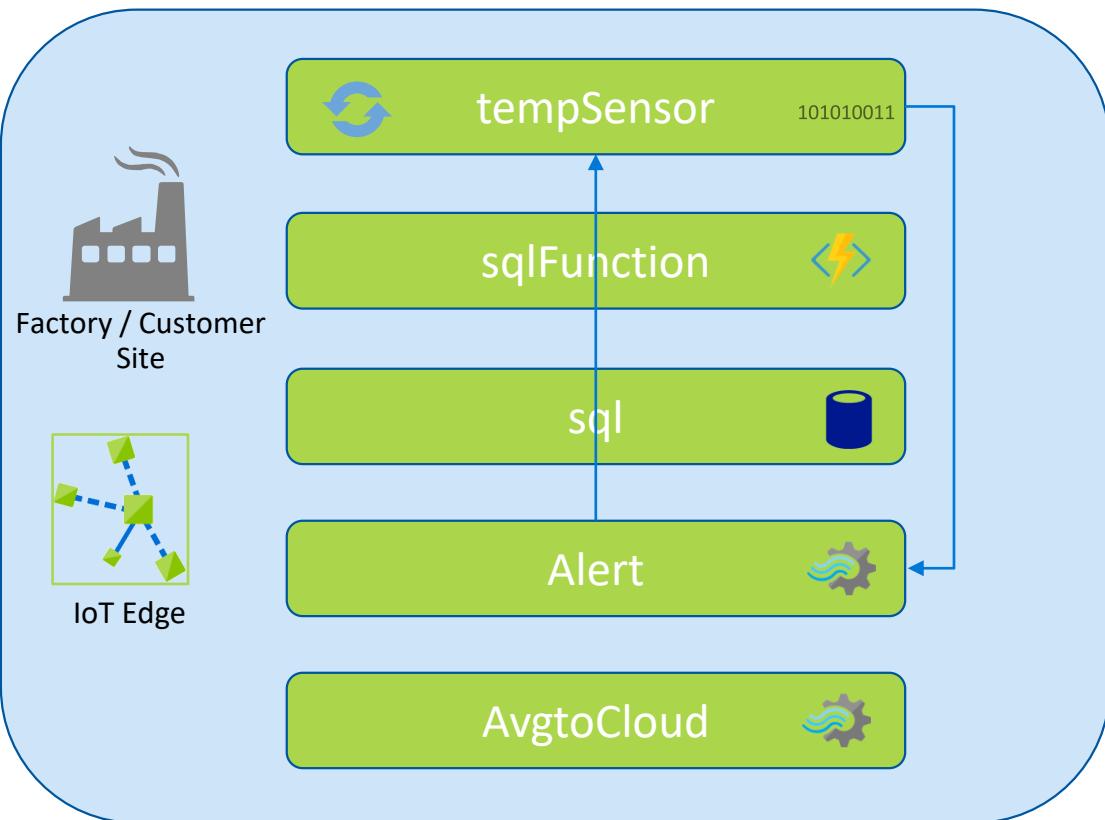
Low Latency

Resiliency

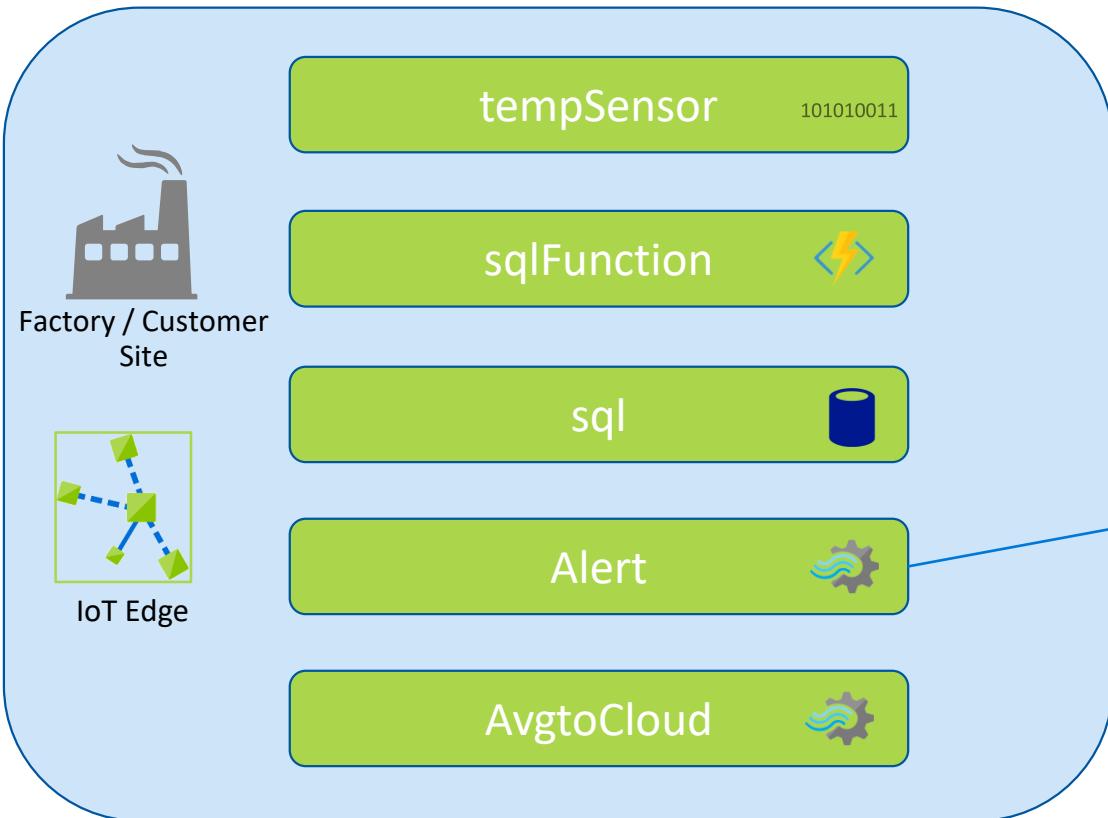
Efficient Use of Bandwidth

Compliance

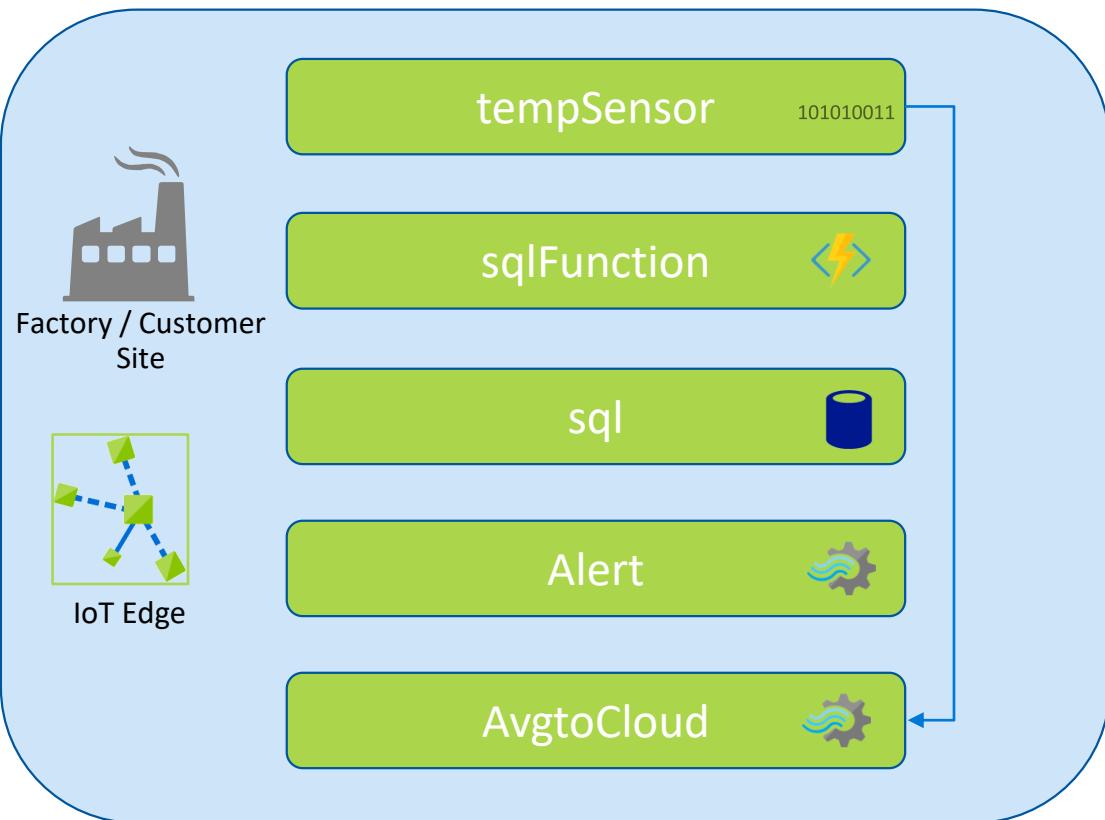
Management at IoT-Scale



```
{
  "routes": {
    "TelemetryToSqlFunction": "FROM /messages/modules/tempSensor/outputs/temperatureOutput INTO BrokeredEndpoint(\"/modules/sqlFunction/inputs/input1\")",
    "TelemetryToAsa": "FROM /messages/modules/tempSensor/* INTO BrokeredEndpoint(\"/modules/Alert/inputs/temperature\")",
    "AlertsToReset": "FROM /messages/modules/Alert/* INTO BrokeredEndpoint(\"/modules/tempSensor/inputs/control\")",
    "AlertsToCloud": "FROM /messages/modules/Alert/* INTO $upstream",
    "TelemetryToAsaAvg": "FROM /messages/modules/tempSensor/* INTO BrokeredEndpoint(\"/modules/AvgtoCloud/inputs/EdgeStream\")",
    "AvgToCloud": "FROM /messages/modules/AvgtoCloud/* INTO $upstream"
  }
}
```



```
{
  "routes": {
    "TelemetryToSqlFunction": "FROM /messages/modules/tempSensor/outputs/temperatureOutput INTO BrokeredEndpoint(\"/modules/sqlFunction/inputs/input1\")",
    "TelemetryToAsa": "FROM /messages/modules/tempSensor/* INTO BrokeredEndpoint(\"/modules/Alert/inputs/temperature\")",
    "AlertsToReset": "FROM /messages/modules/Alert/* INTO BrokeredEndpoint(\"/modules/tempSensor/inputs/control\")",
    "AlertsToCloud": "FROM /messages/modules/Alert/* INTO $upstream",
    "TelemetryToAsaAvg": "FROM /messages/modules/tempSensor/* INTO BrokeredEndpoint(\"/modules/AvgtoCloud/inputs/EdgeStream\")",
    "AvgToCloud": "FROM /messages/modules/AvgtoCloud/* INTO $upstream"
  }
}
```



Inputs
1
EdgeStream

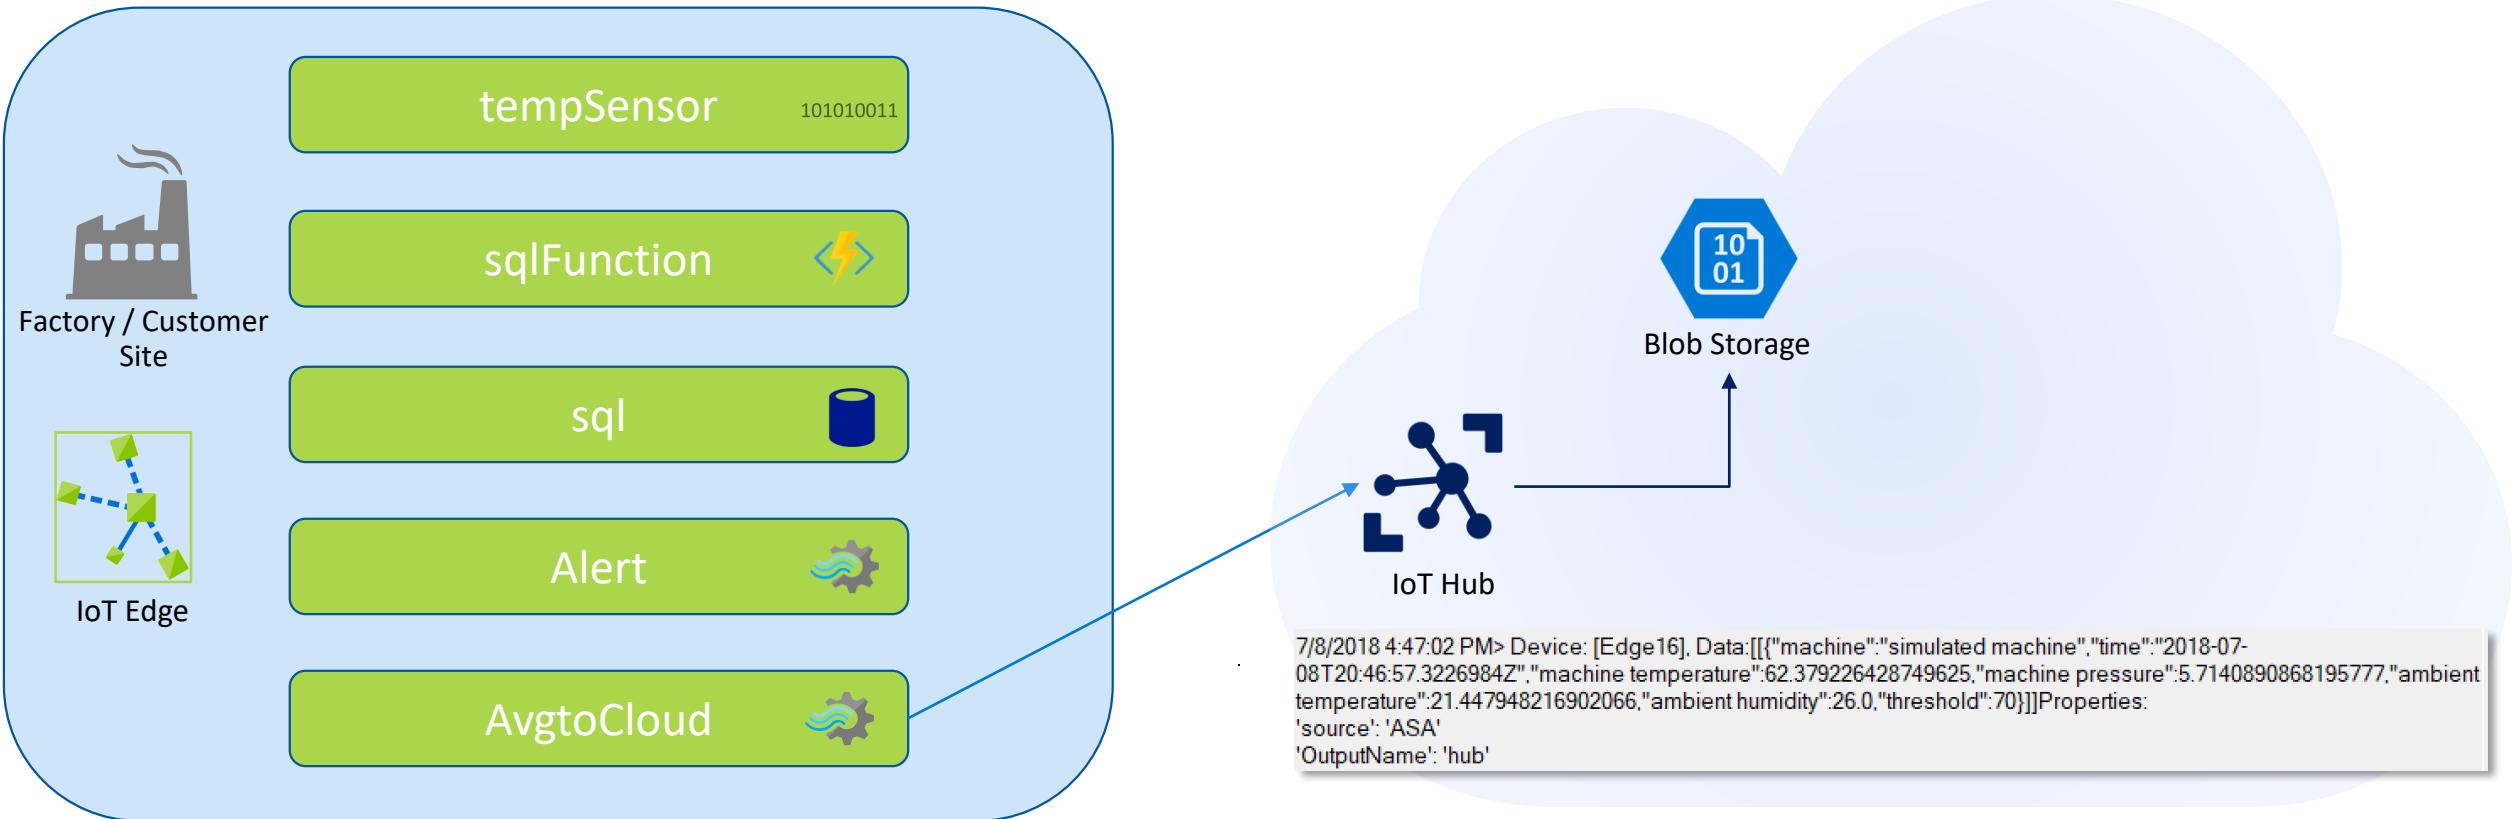
Outputs
1
Hub

```

WITH Telemetry AS (
    SELECT
        'simulated machine' as machine,
        MAX(timeCreated) as time,
        MAX(machine.temperature) as 'machine temperature',
        MAX(machine.pressure) as 'machine pressure',
        MAX(ambient.temperature) as 'ambient temperature',
        MAX(ambient.humidity) as 'ambient humidity',
        70 as threshold
    FROM
        EdgeStream TIMESTAMP BY timeCreated
    WHERE machine.temperature IS NOT NULL
    GROUP BY TumblingWindow(second,30)
)
SELECT
    Tel.*
    INTO Hub
    from Telemetry Tel

```

```
{
    "routes": {
        "TelemetryToSqlFunction": "FROM /messages/modules/tempSensor/outputs/temperatureOutput INTO BrokeredEndpoint(\"/modules/sqlFunction/inputs/input1\")",
        "TelemetryToAsa": "FROM /messages/modules/tempSensor/* INTO BrokeredEndpoint(\"/modules/Alert/inputs/temperature\")",
        "AlertsToReset": "FROM /messages/modules/Alert/* INTO BrokeredEndpoint(\"/modules/tempSensor/inputs/control\")",
        "AlertsToCloud": "FROM /messages/modules/Alert/* INTO $upstream",
        "TelemetryToAsaAvg": "FROM /messages/modules/tempSensor/* INTO BrokeredEndpoint(\"/modules/AvgtoCloud/inputs/EdgeStream\")",
        "AvgToCloud": "FROM /messages/modules/AvgtoCloud/* INTO $upstream"
    }
}
```

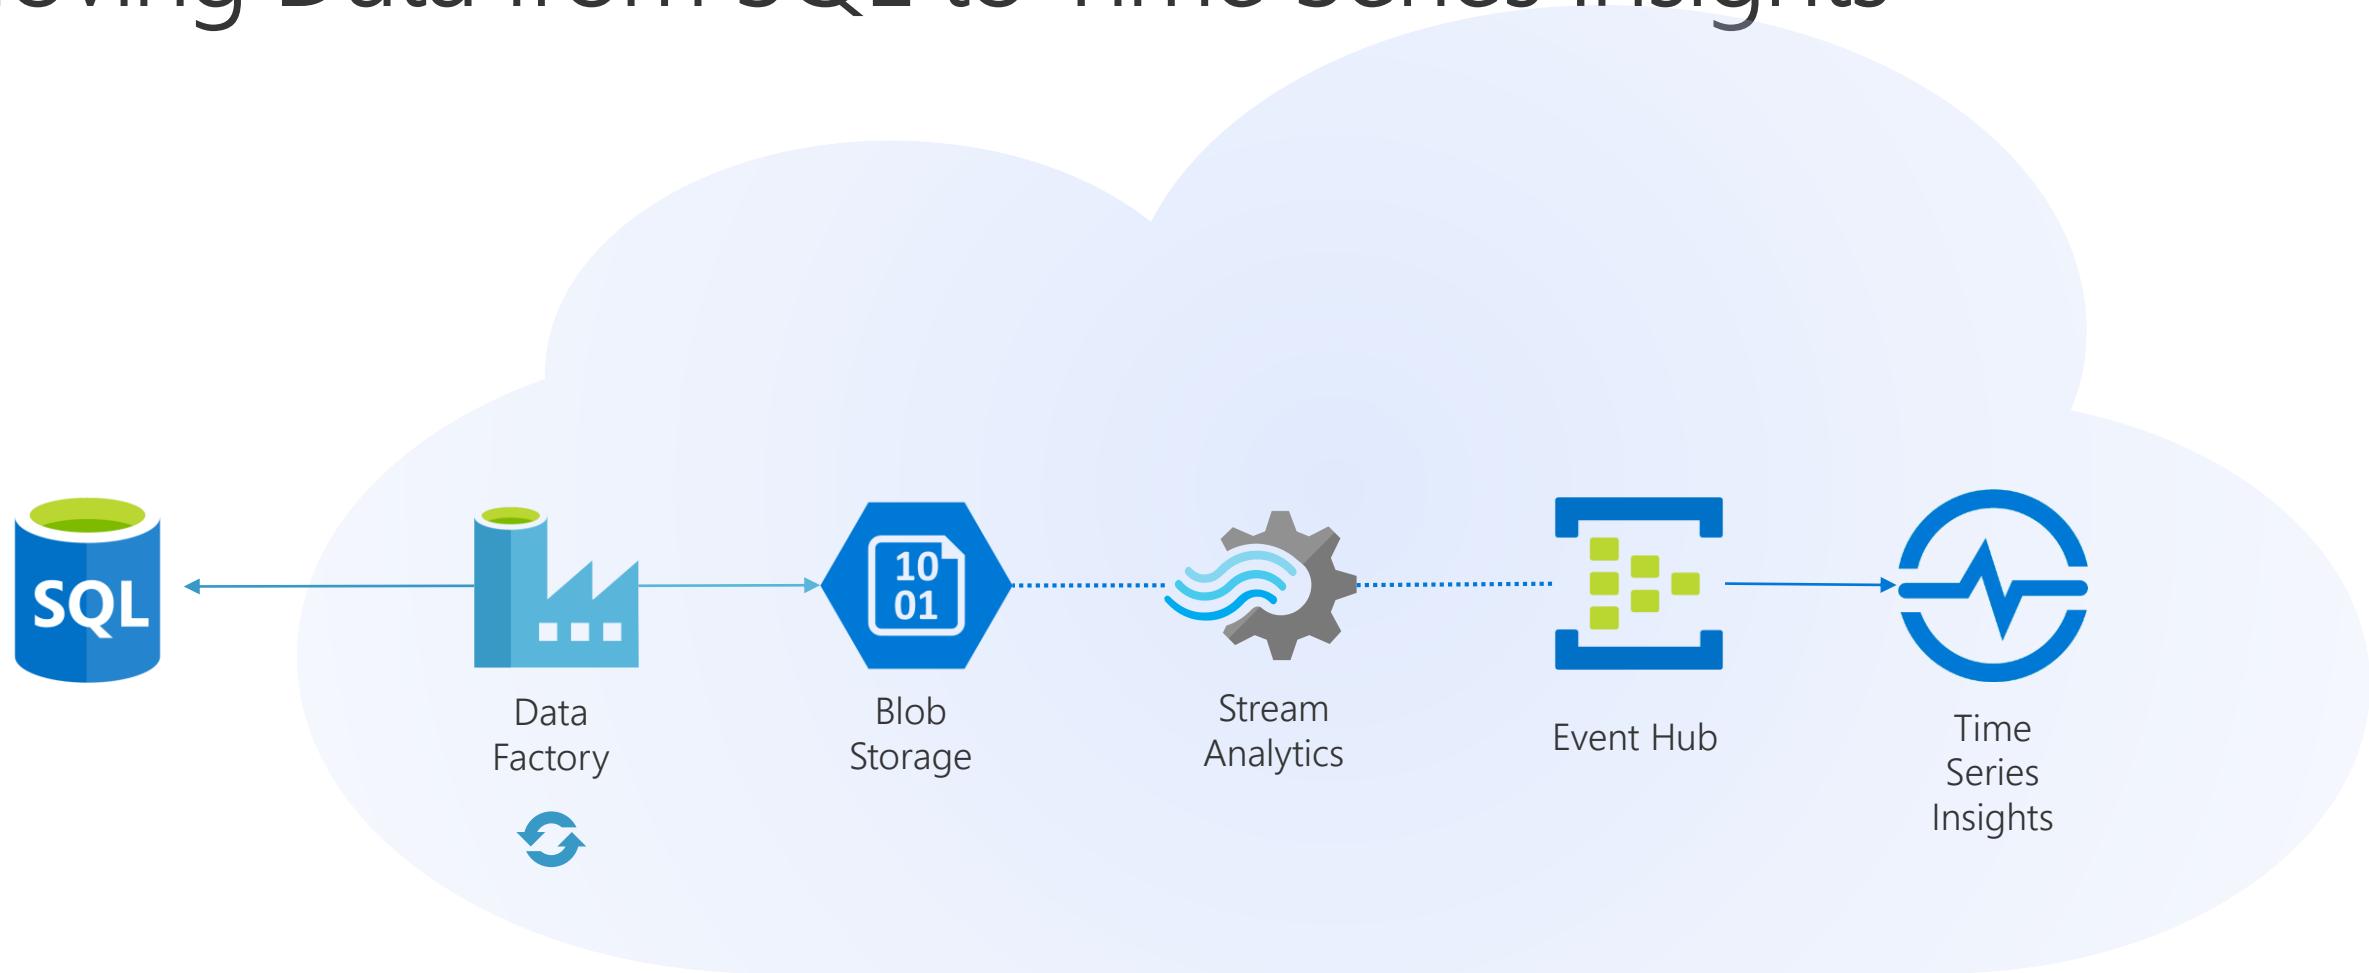


```
{
  "routes": {
    "TelemetryToSqlFunction": "FROM /messages/modules/tempSensor/outputs/temperatureOutput INTO BrokeredEndpoint(\"/modules/sqlFunction/inputs/input1\")",
    "TelemetryToAsa": "FROM /messages/modules/tempSensor/* INTO BrokeredEndpoint(\"/modules/Alert/inputs/temperature\")",
    "AlertsToReset": "FROM /messages/modules/Alert/* INTO BrokeredEndpoint(\"/modules/tempSensor/inputs/control\")",
    "AlertsToCloud": "FROM /messages/modules/Alert/* INTO $upstream",
    "TelemetryToAsaAvg": "FROM /messages/modules/tempSensor/* INTO BrokeredEndpoint(\"/modules/AvgtoCloud/inputs/EdgeStream\")",
    "AvgToCloud": "FROM /messages/modules/AvgtoCloud/* INTO $upstream"
  }
}
```



Demo – Data Factory -> TSI

Moving Data from SQL to Time Series Insights





New Features

Feature availability details

Feature	Status	Remarks
SQL Parallel Write	GA	WW rollout in 1 week
Blob O/P partitioning by custom date-time	Public preview	WW rollout in 1 week
C# UDF on IoT Edge	Public preview	Available now
Live testing in Visual Studio	Public preview	Available now
User defined custom repartition count	Public preview	Available now
New built-in ML models for A/D – Edge and Cloud	Private Preview	Access granted upon sign-up
Custom de-serializers on IoT Edge	Private Preview	Access granted upon sign-up
MSI Authentication for egress to ADLS Gen1	Private Preview	Access granted upon sign-up

Built-in ML models for Anomaly Detection

NEW simpler function signatures for Anomaly Detection

```
AnomalyDetection_SpikeAndDip(  
    <scalar_expression>,  
    <confidence>,  
    <historySize>,  
    <mode>)  
OVER  
    ([PARTITION BY <partition key>]  
     LIMIT DURATION(<unit>, <length>))  
[WHEN  
    boolean_expression])
```

```
AnomalyDetection_ChangePoint(  
    <scalar_expression>,  
    <confidence>,  
    <historySize>)  
OVER  
    ([PARTITION BY <partition key>]  
     LIMIT DURATION(<unit>, <length>))  
[WHEN  
    boolean_expression])
```

Output schema

IsAnomaly: A bigint (0 or 1) indicating if the event was anomalous or not.

Score: A/D score (float) indicating how anomalous an event is.

Local testing with live data in Visual Studio

- Faster iterative testing
- Show results in real time
- View Job metrics
- Time policies support

The screenshot shows the Stream Analytics Local Run Results window and a CMD window. The Stream Analytics window has a toolbar with 'Run Locally' highlighted. The CMD window displays job metrics and logs.

Stream Analytics Local Run Results

TollId	EntryTime	LicensePlate	State	Make	Model	VehicleType
25	8/30/2018 5:22:17 AM	YQB 4701	PA	Toyota	Corolla	1
23	8/30/2018 5:22:18 AM	RFQ 2880	CT	Toyota	Camry	1
5	8/30/2018 5:22:18 AM	ZXC 3799	OR	Ford	Mustang	1
52	8/30/2018 5:22:20 AM	FYW 2785	CT	Ford	Focus	1
18	8/30/2018 5:22:20 AM	JFT 2176	CT	Honda	Civic	1
7	8/30/2018 5:22:21 AM	HYE 3407	CT	Peterbilt	389	2
5	8/30/2018 5:22:21 AM	LAU 3358	CA	Toyota	RAV4	1
64	8/30/2018 5:22:23 AM	VMD 9600	WA	Honda	CRV	1
0	8/30/2018 5:22:23 AM	ZRR 9041	WA	Volvo	S80	1
64	8/30/2018 5:22:23 AM	EUI 9201	PA	Honda	Accord	1
5	8/30/2018 5:22:24 AM	LZW 7265	CT	Chevy	Malibu	1
38	8/30/2018 5:22:26 AM	KBJ 5849	CA	Honda	Civic	1

C:\WINDOWS\system32\CMD.exe

```
9/10/2018 2:11:47 PM : Info : Test Connection for input 'EntryStream'...
9/10/2018 2:11:48 PM : Info : Test Connection for input 'EntryStream' Successfully.
9/10/2018 2:11:56 PM : Info : Job Start Successfully !
9/10/2018 2:12:06 PM
=====
Metrics report=====
DroppedOrAdjustedEvents      :          0
InputEventBytes               :    788.00 KiB
Errors                         :          0
LateInputEvents                :          0
InputEvents                    :      3.53 K
OutputEvents                   :      1.76 K
ConversionErrors              :          0
EarlyInputEvents               :          0
DeserializationError          :          0
=====
9/10/2018 2:12:11 PM
=====
Metrics report=====
DroppedOrAdjustedEvents      :          0
InputEventBytes               :    788.00 KiB
Errors                         :          0
LateInputEvents                :          0
InputEvents                    :      3.53 K
OutputEvents                   :      1.76 K
ConversionErrors              :          0
EarlyInputEvents               :          0
DeserializationError          :          0
=====
```

Custom output partitioning to Blob Storage

Partition egress to Blob storage by

- 1) Any input field
- 2) Custom date and time formats

Gain more fine grain control over data written to Blob storage for dashboarding and reporting

Better alignment with Hive conventions for blob output to be consumed by HDInsight and Azure Databricks.

Path pattern:

`year={datetime:yyyy}/month={datetime:MM}/day={datetime:dd}/hour={datetime:HH}`

Azure blob storage folder nesting example:

`year=2018`

`month=07`

`day=31`

`hour=05`

`hour=06`

The screenshot shows two identical 'Output details' configuration screens for a Stream Analytics job named 'BlobTest'. Both screens are set to 'Provide Blob storage settings manually'. They both point to a 'Storage account' named 'dubansaleastus' and a 'Container' named 'ignitetestoutput'. The 'Path pattern' is highlighted in blue in both cases, showing different values: 'logs/{client_id}/' for configuration 1 and 'cluster1/year={datetime:yyyy}/month={datetime:MM}' for configuration 2. Both configurations also show 'Date format' as 'YYYY/MM/DD', 'Time format' as 'HH', 'Event serialization format' as 'JSON', 'Encoding' as 'UTF-8', and 'Format' as 'Line separated'. A 'Save' button is at the bottom left, and a note about cross-region data movement is at the bottom right.

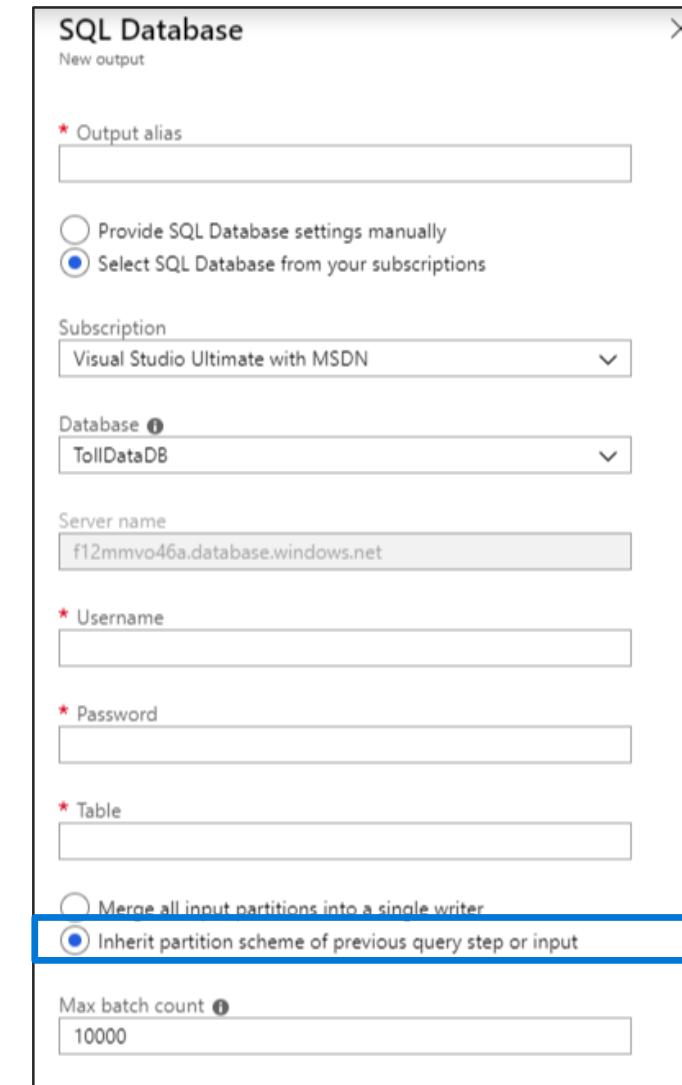
Hi-throughput output to SQL

To achieve fully parallel topologies, ASA will transition SQL 'writes' from Serial to Parallel operations for SQL DB and SQL Data Warehouse

4x-5x improvement in write throughput

Allows for batch size customization to achieve higher throughput

For e.g., this feature enabled our customer building a connected car scenario to scale up from 150K events/min to 500K events/min



ADLS Service Principal Authentication

MSI based authentication will enable egress to Azure Data Lake Storage.

Key benefits over existing AAD (Azure Active Directory) based authentication:

- Job deployment automation (thru Power Shell etc.)
- Long running production jobs
- Consistency with other services

The screenshot shows two side-by-side configuration panels for an Azure Stream Analytics job named "SampleASAWithMsi".

Left Panel: SampleASAWithMsi - Managed service identity

- Header: Home > SampleASAWithMsi - Managed service identity
- Job Name: SampleASAWithMsi - Managed service identity
- Stream Analytics job
- Search (Ctrl+ /)
- Use Managed Service Identity
- Navigation menu:
 - Overview
 - Activity log
 - Access control (IAM)
 - Tags
 - Diagnose and solve problems
- Settings
 - Locks
 - Job topology
 - Inputs
 - Functions
 - Query
 - Outputs
- Configure
 - Scale
 - Locale
 - Event ordering
 - Error policy
 - Compatibility level
 - Managed service identity (highlighted in blue)
- General

Right Panel: Data Lake Storage Gen1

- New output
- * Output alias: DataLakeStorageOutputWithMsi
- Provide Data Lake Storage Gen1 settings manually
- Select Data Lake Storage Gen1 from your subscriptions
- Subscription: Microsoft Azure Internal Consumption (1f824502-de72-4ef...)
- Account name: adla101adls
- * Path prefix pattern: tmp/logs
Example: cluster1/logs/(date)/(time)
- Date format: YYYY/MM/DD
- Time format: HH
- * Event serialization format: JSON
- Encoding: UTF-8
- Format: Line separated
- Authentication mode: Managed service identity

User specified custom re-partition count

Enables better performance tuning

Key Scenarios

- When upstream partition count can't be changed
- Partitioned processing is needed to scale out to larger processing load
- Fixed number of output partitions

```
SELECT *
INTO      [output]
FROM      [input]
PARTITION BY
          DeviceID INTO 10
```

C# UDF support on Azure IoT Edge

Easily extend existing query language with C# User Defined Functions to enable new possibilities, including:

- Complex math calculations
- Machine learning on Azure Edge w/ ML.NET
- String/Date manipulations
- Custom data imputation logic
- Debugging experience in Visual Studio

```
1  using System;
2  using System.Collections.Generic;
3  using System.IO;
4  using System.Linq;
5  using System.Text;
6
7  namespace ASAEdgeUDFDemo
8  {
9      public class Class1
10     {
11         // Add your public static function
12         public static Int64 SquareFunction(Int64 a)
13         {
14             return a * a;
15         }
16     }
17 }
18
```

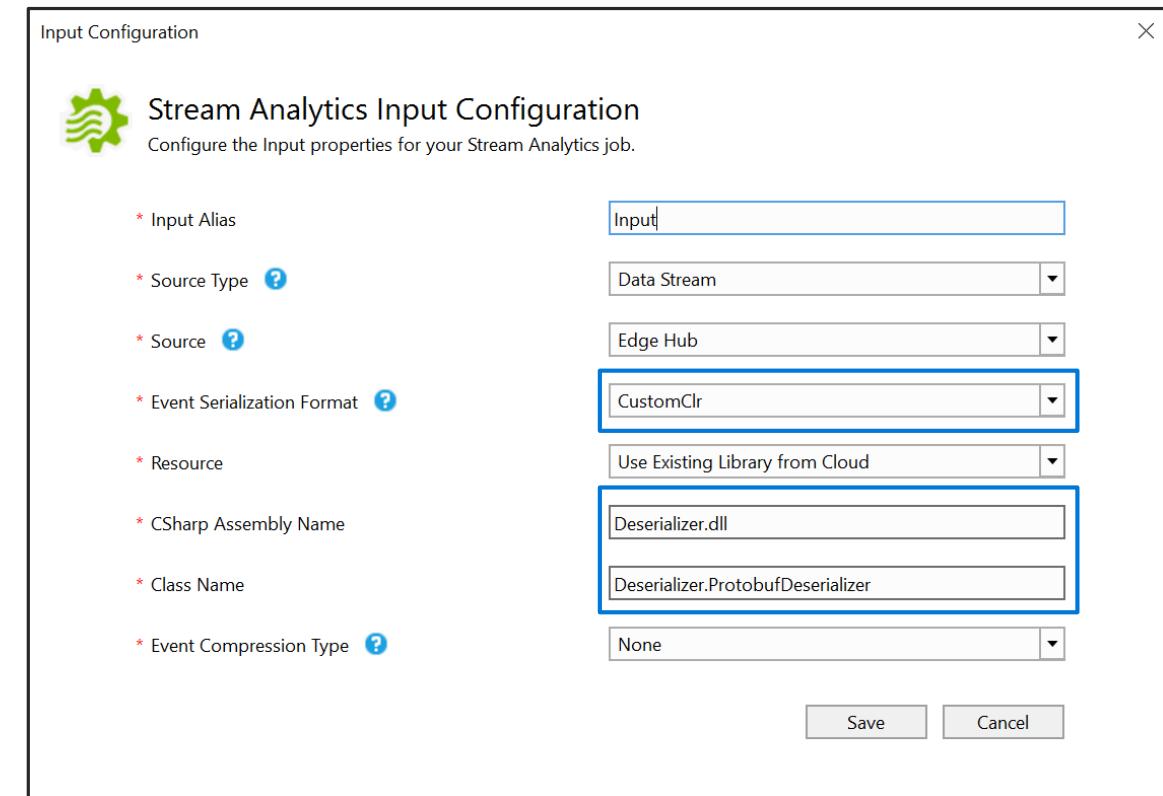
```
▶ Run Locally  Submit to Azure
Select machine.temperature,
udf.ASAEdgeUDFDemo_Class1_SquareFunction(try_cast(machine.tecommperature as bigint))
into Output
from Input
```

Custom de-serializers on Azure IoT Edge

Increases relevancy of Azure Stream Analytics across scores of new scenarios by extending support for data formats beyond JSON, CSV and AVRO

Use custom de-serializers to support a variety of data formats including:

- Protobuf
- Parquet
- XML
- Any binary format!!





Interacting with Stream Analytics

Modalities of Interaction & Programmability

Azure Portal

Visual Studio

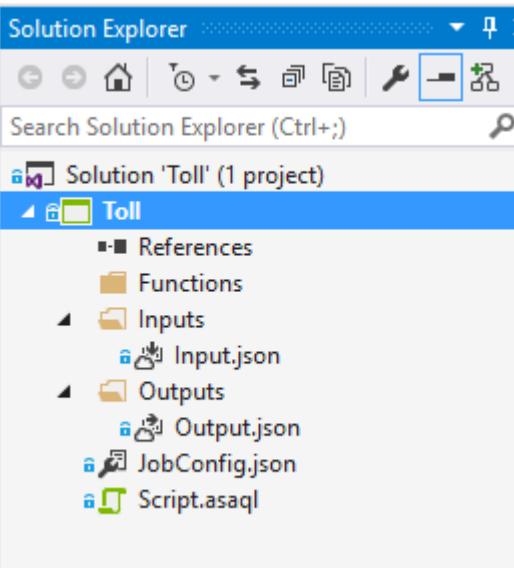
PowerShell

.NET SDK

REST APIs



Visual Studio Support



Version Control
Author with Intellisense
Test Locally
Deploy to Azure
Monitor

The screenshot shows the Azure Stream Analytics Job Overview page within Visual Studio. It includes the following sections:

- New Project**: Shows recent and installed projects, with 'Azure Stream Analytics Application' highlighted.
- Job Summary**: Displays job status as 'Running', created time as '12/12/2016 4:47:05 PM', and last output time as '12/13/2016 11:06:00 AM'.
- Job Metrics (Last 30 Mins)**: Provides real-time metrics:

Total Input Events	4.95 K
Input Event Bytes	946.06 KiB
Total Output Events	896
Late Input Events	11
Out of Order Events	0
Function Events	0
Function Requests	0
Failed Function Requests	0
Data Conversion Errors	0
Runtime Errors	0
- Script.asaql**: A preview window showing the ASA query:

```
SELECT TollId, System.Timestamp
FROM EntryStream TIMESTAMP BY EntryTime
GROUP BY TUMBLINGWINDOW(minute, 5)
```

with options to 'Run Locally' or 'Submit To Azure'.



Q & A

Thank You!

ευχαριστώ

Salamat Po

متشکرم

شكراً

Grazie

благодаря

ありがとうございます

Kiitos Teşekkürler 谢谢

ឧបម្ពុណ្ឋរំ

Obrigado

شكريه

Terima Kasih

Dziękuję

Hvala

Köszönöm

Tak

Dank u Wel

дякую

Tack

Mulțumesc

спасибо

Danke

Cám ơn

Gracias

多謝晒

Ďakujem

הודות

ନଞ୍ଚି

Děkuji

감사합니다



© 2018 Microsoft Corporation. All rights reserved. Microsoft, Windows, and other product names are or may be registered trademarks and/or trademarks in the U.S. and/or other countries. The information herein is for informational purposes only and represents the current view of Microsoft Corporation as of the date of this presentation. Because Microsoft must respond to changing market conditions, it should not be interpreted to be a commitment on the part of Microsoft, and Microsoft cannot guarantee the accuracy of any information provided after the date of this presentation. MICROSOFT MAKES NO WARRANTIES, EXPRESS, IMPLIED OR STATUTORY, AS TO THE INFORMATION IN THIS PRESENTATION.